

Statistical Learning

Common Core

Understanding Machine Learning

General Recipe

statistical learning algorithm by combining:

- **model** (e.g., linear function, Gaussian distribution)
- **objective function** (e.g., squared residuals)
- **optimization algorithm** (e.g., gradient descent)

Consider Supervised Learning Scenario

map input to output: $y = f(\mathbf{x})$ (estimated: $\hat{f}(\mathbf{x})$)

random variables Y and $\mathbf{X} = (X_1, X_2, \dots, X_p) \leftarrow$ usually high-dimensional

curve fitting / parameter estimation:

fit train data set of (y_i, \mathbf{x}_i) pairs \rightarrow minimization of cost function

consider discriminative models:

- predict conditional density function $p(y|\mathbf{x}_i)$
(as opposed to generative models predicting $p(y, \mathbf{x}) \rightarrow \mathbf{x}$ not given, more difficult)
- often just conditional mean (depending on used loss) $E[Y|\mathbf{X} = \mathbf{x}_i]$ of $p(y|\mathbf{x}_i)$

Model

Model in Linear Regression

fit:

$$y_i = \hat{\alpha} + \overbrace{\sum_{j=1}^p \hat{\beta}_j x_{ij}}^{\hat{f}(\mathbf{x}_i)} + \varepsilon_i$$

to be estimated:

- $\hat{\alpha}, \hat{\beta}$

$$\rightarrow \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n \left(y_i - \hat{f}(\mathbf{x}_i) \right)^2$$

(approximating assumed true α, β, σ)

predict:

$$\hat{y}_i = E[Y|\mathbf{X} = \mathbf{x}_i] = \hat{f}(\mathbf{x}_i)$$

$$p(y|\mathbf{x}_i) = \mathcal{N}(y; \hat{y}_i, \hat{\sigma}^2)$$

Gaussian

mean

variance
(reflected
by ε_i in fit)

Model Ingredients

- underlying probability distribution of Y
- approximated functional form $\hat{f}(\mathbf{x}) \rightarrow$ depends on used ML algorithm

$$p(y|\mathbf{x}_i) = \text{PDF}(y; \hat{f}(\mathbf{x}_i), \dots)$$

assumed Probability
Density Function of target

$\hat{f}(\mathbf{x})$ often estimation
of location parameter

other parameters
(e.g., scale parameter)
often not predicted
per sample (assume
homoscedasticity)

Principle of Maximum Entropy

How to choose the right model, i.e., the assumed underlying probability distribution of Y generating the observed data?

Occam's razor: pick distribution with fewest assumptions

→ distribution with maximum entropy under given range and moment constraints

discrete: $H(X) = -\sum p_k \log p_k$, continuous: $H(X) = -\int p(x) \log p(x) dx$

(→ bulk of entropy from tails of distributions)

examples for maximum entropy distributions:

- normal distribution: continuous on $(-\infty, \infty)$ with particular mean and variance
- Poisson distribution: discrete on $[0, \infty)$ with particular mean

Objective Function

Loss Function

loss function L : expressing deviation between prediction and target

$$L(y_i, \hat{f}(\mathbf{x}_i); \hat{\boldsymbol{\theta}})$$

with $\hat{\boldsymbol{\theta}}$ corresponding to parameters of model $\hat{f}(\mathbf{x})$

e.g., $\hat{\alpha}, \hat{\boldsymbol{\beta}}$ in linear regression

e.g., squared residuals (for regression problems):

$$L(y_i, \hat{f}(\mathbf{x}_i); \hat{\boldsymbol{\theta}}) = \left(y_i - \hat{f}(\mathbf{x}_i; \hat{\boldsymbol{\theta}}) \right)^2$$

Cost Function

averaging losses over (empirical) training data set:

$$J(\hat{\theta}) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(\mathbf{x}_i); \hat{\theta})$$

cost function to be minimized according to model parameters $\hat{\theta}$

→ objective function

Optimization

Cost Minimization

minimize training costs $J(\hat{\boldsymbol{\theta}})$ according to model parameters $\hat{\boldsymbol{\theta}}$:

$$\nabla_{\hat{\boldsymbol{\theta}}} J(\hat{\boldsymbol{\theta}}) = 0$$

for mean squared error (aka least squares method):

$$\nabla_{\hat{\boldsymbol{\theta}}} \frac{1}{n} \sum_{i=1}^n \left(y_i - \hat{f}(\mathbf{x}_i; \hat{\boldsymbol{\theta}}) \right)^2 = 0$$

Ordinary Least Squares

linear regression: $\hat{f}(\mathbf{x}_i) = \hat{\alpha} + \sum_{j=1}^p \hat{\beta}_j x_{ij}$

matrix notation (drop intercept α for simplicity here): $\hat{\mathbf{y}} = \mathbf{X} \hat{\boldsymbol{\beta}}$

solve normal equations for training data set:

$$\nabla_{\hat{\boldsymbol{\beta}}} (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}})^T (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}}) = 0$$

$$\rightarrow \hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

but usually no closed-form solution \rightarrow need for numerical optimization

Maximum Likelihood Estimation

Likelihood

special objective function: likelihood function \mathcal{L}

- value of predicted probability density function (model) at observed target
- as function of model parameters $\hat{\theta}$

$$\mathcal{L}(\hat{\theta}; \mathbf{y}|\mathbf{x}) = p(\mathbf{y}|\mathbf{x}; \hat{\theta})$$

a little bit confusing:

$p(\mathbf{y}|\mathbf{x}; \hat{\theta})$ is conditional density function if function of data with $\hat{\theta}$ fixed

$p(\mathbf{y}|\mathbf{x}; \hat{\theta})$ is conditional likelihood function if function of $\hat{\theta}$ with data fixed

Independence Assumption

consider training data set (of size n) as one sample from unknown joint probability distribution of n independent (i.i.d.) sets of random variables $(Y_1, \mathbf{X}_1), (Y_2, \mathbf{X}_2), \dots, (Y_n, \mathbf{X}_n)$
(same thing as random sampling n times from (Y, \mathbf{X}))

→ likelihood as product of univariate probability density functions:

$$\mathcal{L}(\hat{\boldsymbol{\theta}}; \mathbf{y}|\mathbf{x}) = \prod_{i=1}^n p(y_i|\mathbf{x}_i; \hat{\boldsymbol{\theta}})$$

... to be maximized according to estimated model parameters $\hat{\boldsymbol{\theta}}$
(known as maximum likelihood estimation)

Negative Log-Likelihood

model, e.g., for linear regression:

$$p(y_i|x_i) = \mathcal{N}(y; \hat{\alpha} + x_i \hat{\beta}, \sigma^2)$$

- with \mathbf{y} , \mathbf{X} given in training
- $\hat{\boldsymbol{\theta}} = (\hat{\alpha}, \hat{\beta}, \hat{\sigma}^2)$ to be estimated

logarithmic transformation of \mathcal{L} : log-likelihood function

$$\ell(\hat{\boldsymbol{\theta}}; \mathbf{y}|\mathbf{x}) = \ln(\mathcal{L}(\hat{\boldsymbol{\theta}}; \mathbf{y}|\mathbf{x})) = \sum_{i=1}^n \ln(p(y_i|x_i; \hat{\boldsymbol{\theta}}))$$

- logarithm monotonic function \rightarrow maximum of ℓ and \mathcal{L} at same $\hat{\boldsymbol{\theta}}$ values
- most probability distributions (e.g., exponential family) only logarithmically concave (important for optimization)
- sum computationally more convenient than product

negative log-likelihood: minimization instead of maximization

$$\hat{\boldsymbol{\theta}} = \operatorname{argmax}_{\hat{\boldsymbol{\theta}}} \ell(\hat{\boldsymbol{\theta}}; \mathbf{y}|\mathbf{x}) = \operatorname{argmin}_{\hat{\boldsymbol{\theta}}} \underbrace{(-\ell(\hat{\boldsymbol{\theta}}; \mathbf{y}|\mathbf{x}))}_{\text{objective function}}$$

Maximum Likelihood

estimate parameters $\hat{\boldsymbol{\theta}} = (\hat{\theta}_1, \hat{\theta}_2, \dots, \hat{\theta}_k)$ solving

$$\hat{\boldsymbol{\theta}} = \operatorname{argmax}_{\hat{\boldsymbol{\theta}}} \ell(\hat{\boldsymbol{\theta}}; \mathbf{y} | \mathbf{x})$$

mode of likelihood
as point estimate

→ likelihood equations (if ℓ differentiable in $\hat{\boldsymbol{\theta}}$):

$$\frac{\partial \ell}{\partial \hat{\theta}_1} = 0, \frac{\partial \ell}{\partial \hat{\theta}_2} = 0, \dots, \frac{\partial \ell}{\partial \hat{\theta}_k} = 0 \quad (\text{in short: } \nabla_{\hat{\boldsymbol{\theta}}} \ell = 0)$$

can be solved explicitly for some cases (e.g., ordinary least squares for linear regression)

but usually no closed-form solution

→ need for numerical optimization (e.g., gradient descent)

least squares method

corresponds to maximum likelihood estimation with

Gaussian model: $\sim e^{-\frac{1}{2}\left(\frac{y-\hat{\mu}}{\hat{\sigma}}\right)^2}$

Interpretation of Maximum Likelihood

quantification of
difference between two
probability distributions

maximum likelihood estimation corresponds to minimization of Kullback-Leibler divergence (as well as cross entropy: $D_{KL}(p||q) = H(p, q) - H(p)$) between true data-generating probability distribution (manifested by empirical distribution of training data) and probability distribution of model:

$$\begin{aligned} \operatorname{argmin}_{\hat{\theta}} D_{KL} \left(p_{\text{model}}(y|\mathbf{x}; \hat{\theta}) || p_{\text{data}}(y|\mathbf{x}) \right) &= \operatorname{argmin}_{\hat{\theta}} \int p_{\text{data}}(y|\mathbf{x}) \log \frac{p_{\text{data}}(y|\mathbf{x})}{p_{\text{model}}(y|\mathbf{x}; \hat{\theta})} dy \\ &= \operatorname{argmin}_{\hat{\theta}} \frac{1}{n} \sum_{i=1}^n \log \frac{p_{\text{data}}(y_i|\mathbf{x}_i)}{p_{\text{model}}(y_i|\mathbf{x}_i; \hat{\theta})} = \operatorname{argmax}_{\hat{\theta}} \sum_{i=1}^n \ln \left(p_{\text{model}}(y_i|\mathbf{x}_i; \hat{\theta}) \right) = \operatorname{argmax}_{\hat{\theta}} \ell(\hat{\theta}; \mathbf{y} | \mathbf{x}) \end{aligned}$$

$E[\dots]$ no contribution from p_{data}

make the model distribution match the empirical distribution

(mean squared error corresponds to cross-entropy between empirical distribution and Gaussian model)

Iterative Optimization

Gradient Descent

usually (except for special cases like ordinary least squares) no closed-form solution to ML optimization problems like minimization of a cost function or maximization of a likelihood function:

$$\nabla_{\hat{\theta}} J(\hat{\theta}) = 0$$

→ need for numerical methods

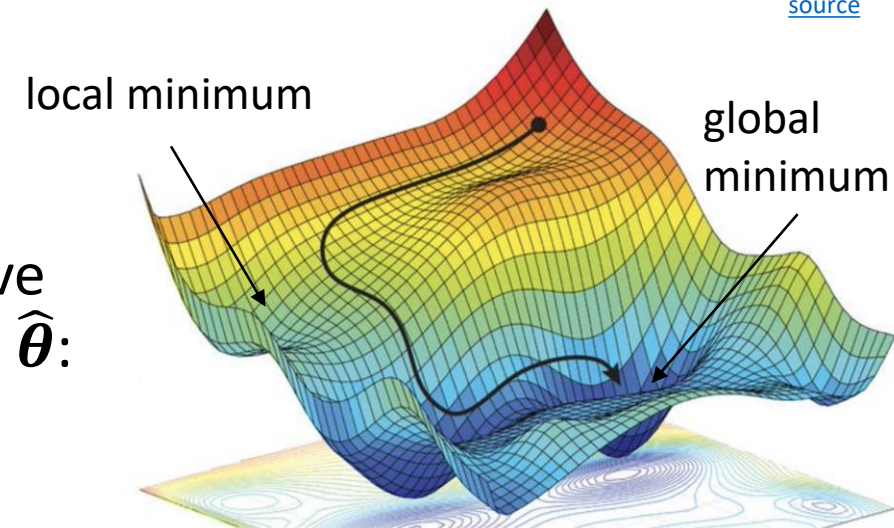
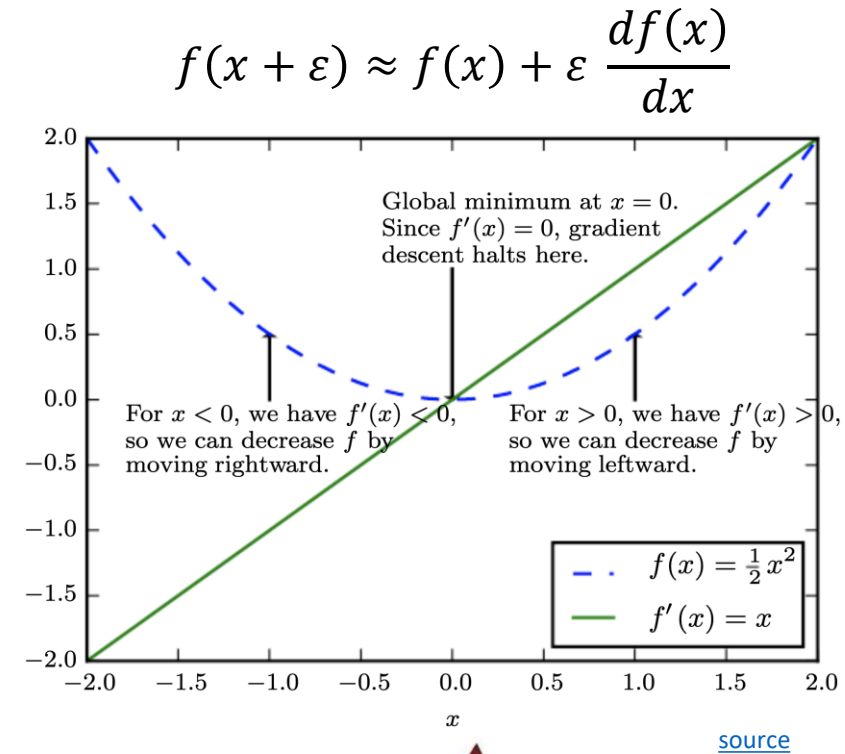
most popular choice: gradient descent

decreasing J by iteratively moving in direction of negative gradient (steepest descent) with respect to input vector $\hat{\theta}$:

$$\hat{\theta} \leftarrow \hat{\theta} - \eta \nabla_{\hat{\theta}} J(\hat{\theta})$$

step size
(learning rate)

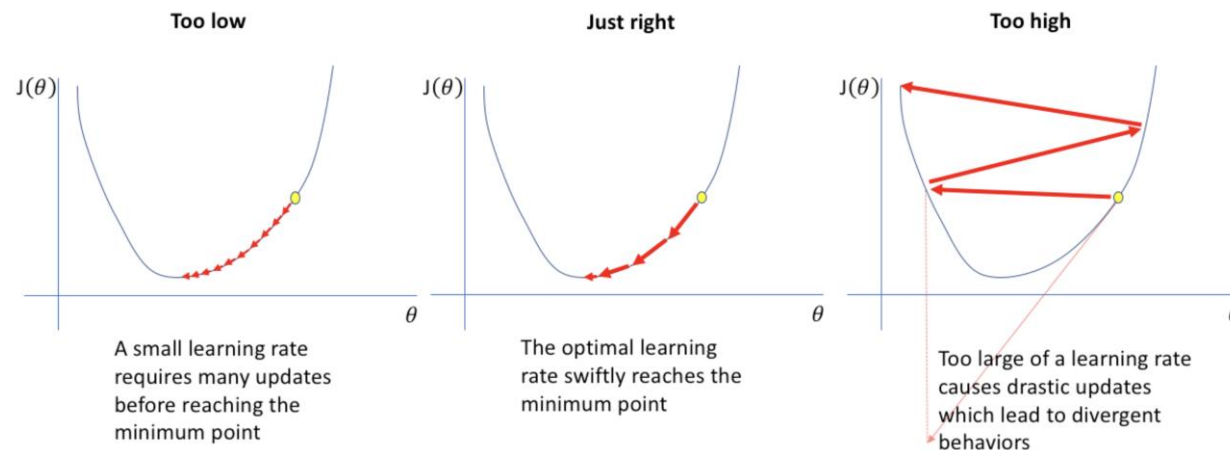
vector containing all partial derivatives



Learning Rate

learning rate η corresponds to positive scalar value (hyperparameter)

- often set to small constant
- can be optimized via line search: evaluate $J\left(\hat{\theta} - \eta \nabla_{\hat{\theta}} J(\hat{\theta})\right)$ for several values of η and choose η resulting in smallest value of objective function
- can be varied from iteration to iteration (e.g., via some heuristic): smaller steps closer to the minimum (learning rate schedule/decay)



Batch and Stochastic Gradient Descent

in usual gradient descent: parameters $\hat{\theta}$ (and in turn objective function J) updated after full training epoch (one sweep through entire training data set) \rightarrow batch learning

consider: $J(\hat{\theta}) = \frac{1}{n} \sum_{i=1}^n J_i(\hat{\theta})$ (J_i corresponds to loss function)

stochastic gradient descent (SGD) updates after each training example (gradient of $J(\hat{\theta})$ approximated with gradient of single loss): $\hat{\theta} \leftarrow \hat{\theta} - \eta \nabla_{\hat{\theta}} J_i(\hat{\theta})$

\rightarrow online learning

- shuffling of samples after training epoch to prevent cycles
- adaptive learning rate to improve convergence

Mini Batches

Compared to its batch mode, stochastic gradient descent helps to avoid local minima (more robust) but is computationally expensive (less efficient).

→ mini-batch SGD as compromise:

splitting training data set into small batches used to calculate model error and update parameters

at the cost of one additional hyperparameter specifying mini-batch size


Gradient Descent with Momentum

goal: improve optimization process by avoiding gradients bouncing around the search space (dampening oscillations), accelerate in direction of minima

algorithm:

- estimate gradient \mathbf{g}
- compute velocity update: $\mathbf{v} \leftarrow \alpha \mathbf{v} - \eta \mathbf{g}$
- apply parameter update: $\hat{\boldsymbol{\theta}} \leftarrow \hat{\boldsymbol{\theta}} + \mathbf{v}$

hyperparameter ($0 < \alpha < 1$) specifying exponential decay (like learning rate η , α may be adapted over course of optimization)



velocity: direction and speed at which the parameters move through parameter space, set to an exponentially decaying average of the past negative gradients (continues to move in its direction)

Bayesian Methods

Bayesian Statistics

frequentist perspective:

- true model parameters θ fixed (but unknown \rightarrow uncertain point estimates $\hat{\theta}$)
- data set from random sampling

Bayesian perspective:

- true model parameters θ random (estimation of full distributions over parameters)
- data not random (observed)

Bayes Theorem

H : hypothesis (model)

E : evidence (data)

likelihood function (derived
from statistical model)

$$P(H|E) = \frac{P(E|H) \cdot P(H)}{P(E)}$$

posterior probability
(belief updated with
observed evidence)

new data (same
for all hypotheses
→ uninteresting)

prior probability
(belief in H
before evidence)

Bayes Theorem for Parameter Estimation

data likelihood:
model (function of θ with data fixed), e.g.,
for linear regression: $\mathcal{N}(y; \alpha + \mathbf{x} \boldsymbol{\beta}, \sigma^2)$

$$p(\theta|y, \mathbf{x}) \propto p(y|\mathbf{x}, \theta) \cdot p(\theta)$$

posterior:
probability distribution reflecting
the effect of data on prior belief
about parameters (increasing
certainty with new evidence)

to be compared
to frequentist
point estimates $\hat{\theta}$

prior (e.g., broad Gaussian):
probability distribution reflecting
initial belief about uncertain
parameters (helps to reduce
variance with few data)

Bayesian Prediction

$$p(y_{n+1} | \mathbf{x}_{n+1}, (y_1, \mathbf{x}_1), \dots, (y_n, \mathbf{x}_n)) = \int \underbrace{p(y_{n+1} | \mathbf{x}_{n+1}, \boldsymbol{\theta})}_{\text{likelihood}} \underbrace{p(\boldsymbol{\theta} | (y_1, \mathbf{x}_1), \dots, (y_n, \mathbf{x}_n))}_{\text{latest posterior}} d\boldsymbol{\theta}$$

prediction of full probability distribution instead of frequentist point estimate(s) plugged into model distribution ($p(y|\mathbf{x}) = \text{PDF}(y; \hat{f}(\mathbf{x}), \dots)$)

drawback: high computational costs due to integral over parameter distribution (especially for large training data sets)

Alternatives for PDF Predictions

quantile regression: estimate quantile τ of distribution instead of conditional mean by minimizing pinball loss

$$(1 - \tau) \sum_{y_i < \hat{q}_i} (\hat{q}_i - y_i) + \tau \sum_{y_i \geq \hat{q}_i} (y_i - \hat{q}_i)$$

instead of squared error loss (choice of loss function defines point estimate)

- possible with various ML methods, including neural networks and tree-based methods (like random forests or gradient boosting)
- subsequently approximate full probability distribution from a few predicted quantiles by fitting (assumed distribution or spline) or [quantile-parameterized distributions](#)

PDF assumption for model and estimation of its parameters: assume, e.g., Gaussian or negative binomial and estimate mean and variance (by means of different ML models or corresponding likelihood)

Common Core of Statistical Learning

Most ML algorithms can be described by the general recipe of combining models, costs, and optimization methods.

including non-linear models:

- neural networks: backpropagation
- support-vector machines: minimizing hinge loss (soft-margin SVM)
- decision trees: minimizing impurity functions (mean squared error for regression, (kind of) Kullback-Leibler divergence for classification)

and even

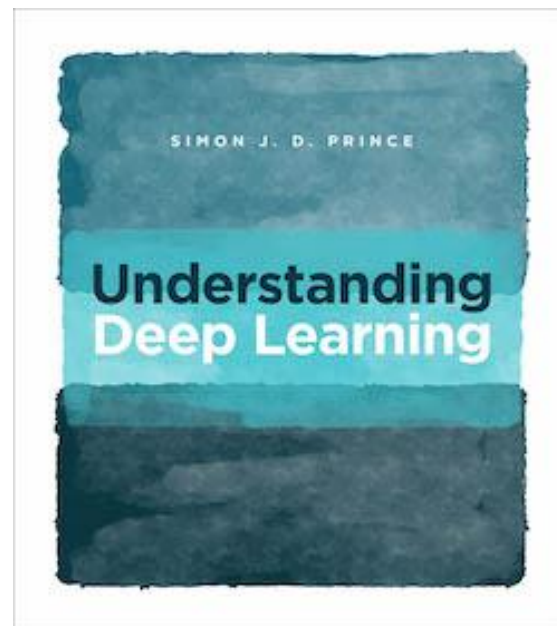
- unsupervised learning: self-supervised, PCA by maximum variance directions
- reinforcement learning: express rewards in loss functions

DEEP LEARNING

Ian Goodfellow, Yoshua Bengio,
and Aaron Courville

Literature

not only deep learning, but also a nice ML
introduction: <https://www.deeplearningbook.org/>



following on from the above,
also covering newer topics:
[udlbook](#)

Overcome our Mathematical Limitations

evolution provided us with moderate math skills

AI/ML to the rescue:

- [recognition of mathematical structures and patterns](#)
- algorithm discovery: [AlphaTensor](#)
- symbolic regression ([e.g., for physics](#))