

Reinforcement Learning

Understanding Machine Learning

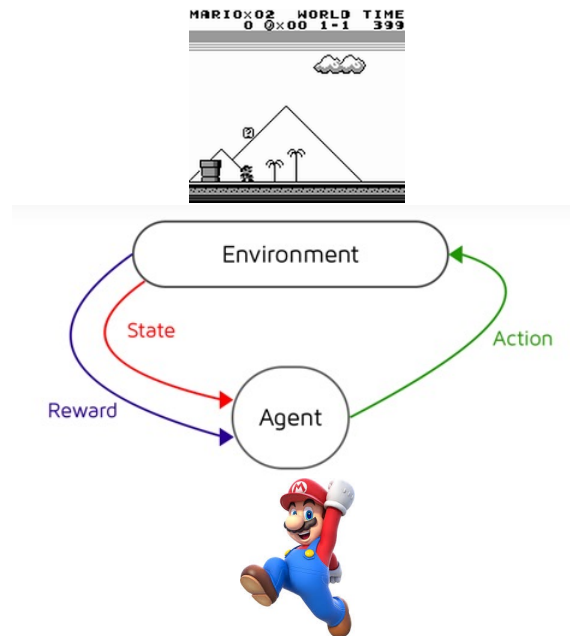
October 2022

Felix Wick

Sequential Decision Making

reinforcement learning (RL):

formalization of sequential decision making of software agent interacting with environment



Main Elements of RL

goal: find action policy maximizing reward from environment

action policy: exploration-exploitation trade-off

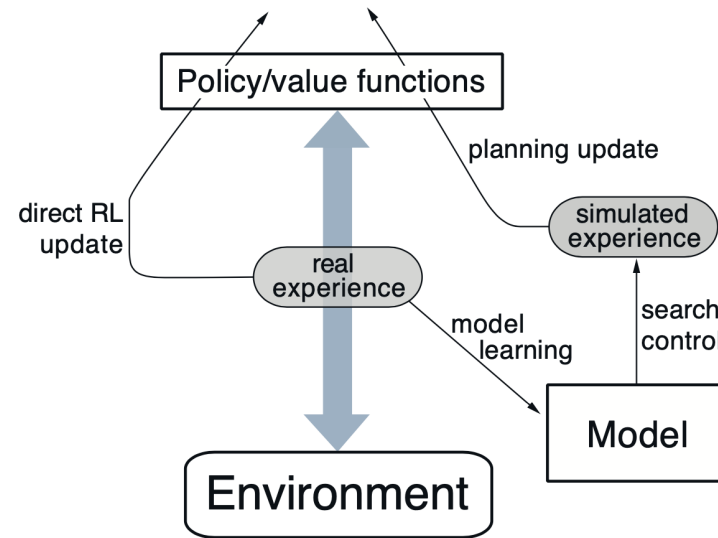
- e.g., epsilon-greedy: random exploration at small fraction of the time
- off-policy instead of on-policy learning: policy for learning different from current best → exploit in application and explore during learning

feedback from environment: goal-directed, no supervision

- scalar reward signal
- cumulative and delayed rewards (credit assignment problem)

Optional Elements of RL

model of environment: (model-free) trial-and-error or planning



from Sutton

value functions for states or actions: improve efficiency of search in vast action policy space (alternative: direct policy search)

Markov Decision Process (MDP)


idea: current state includes all information about past

transition probabilities between states describe dynamics of given MDP

action policy: mapping from states to probabilities for selection of different actions

States, Actions, and Rewards

transition probabilities (model of environment): $p(s_{t+1}, r_{t+1} | s_t, a_t)$



reward hypothesis:

- reward as scalar signal
- goal: maximization of expected cumulative sum of received rewards

Value-Based Methods

State and Action Values

state/action value: total amount of expected future reward starting from given state/action (usually with discounting of later steps)

→ indicating long-term desirability of states/actions

main motivation: improve efficiency of search in policy space

(for comparison: evolutionary methods search directly by evaluating entire policies)

State-Value Function

$$\begin{aligned} v_{\pi}(s_t) &= E_{\pi} \left[\overbrace{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1}}^{\text{return}} \mid s_t \right] = E_{\pi} [r_{t+1} + \underbrace{\gamma v_{\pi}(s_{t+1})}_{\text{discount rate}} \mid s_t] \\ &= \sum_{a_t} \pi(a_t \mid s_t) \sum_{s'_{t+1}, r_{t+1}} p(s'_{t+1}, r_{t+1} \mid s_t, a_t) [r_{t+1} + \gamma v_{\pi}(s'_{t+1})] \end{aligned}$$

policy: probability to take specific action being in a given state

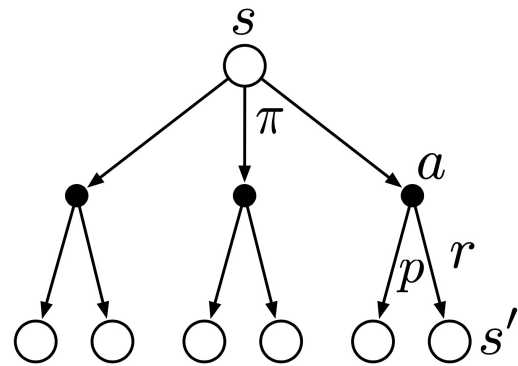
transition probability (depending on environment) from one state to another for a given action

Bellman (expectation) equation: recursion

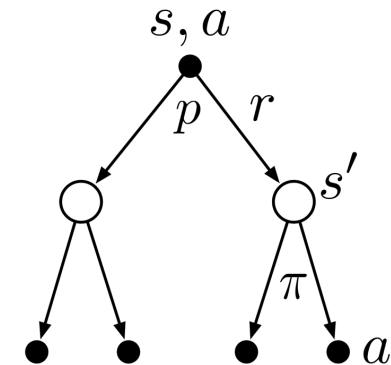
Action-Value Function

$$q_{\pi}(s_t, a_t) = E_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t, a_t \right] = E_{\pi} [r_{t+1} + \gamma q_{\pi}(s_{t+1}, a_{t+1}) \mid s_t, a_t]$$

$$= \sum_{s'_{t+1}, r_{t+1}} p(s'_{t+1}, r_{t+1} \mid s_t, a_t) [r_{t+1} + \gamma v_{\pi}(s'_{t+1})]$$



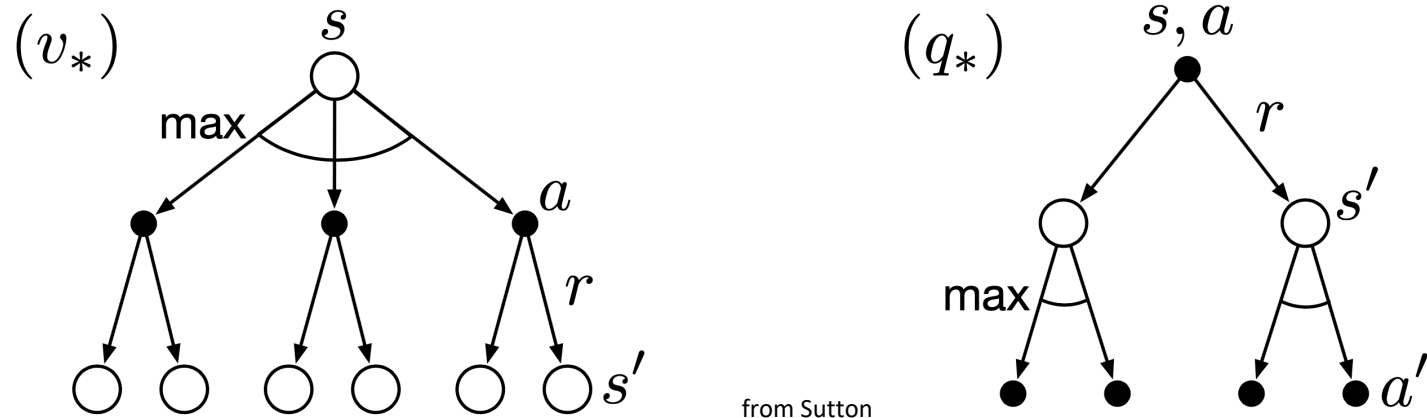
Backup diagram for v_{π}



q_{π} backup diagram

Bellman Optimality Equations

optimal solutions to Bellman equations (directly defining optimal policy):



rarely possible to find in practice (model of environment, Markov property, computational resources)

→ approximate solutions

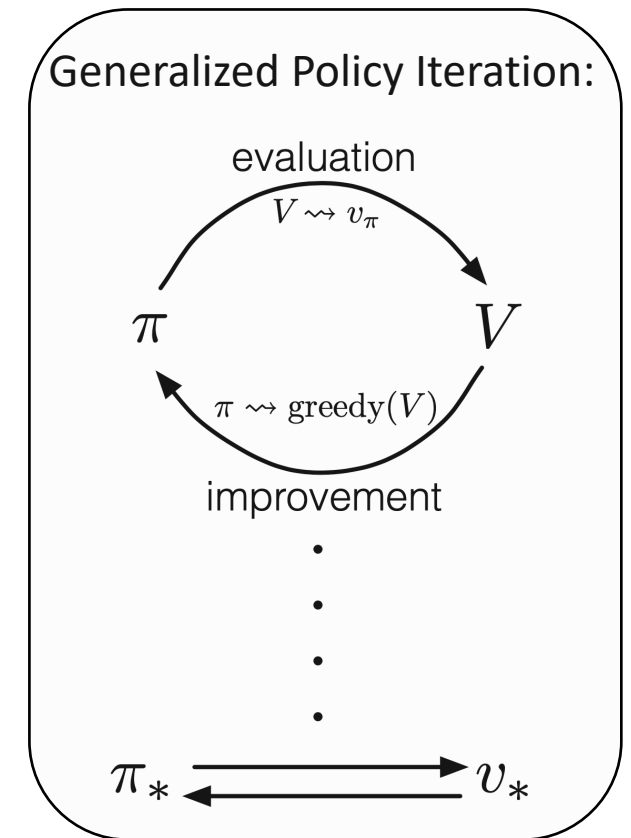
Dynamic Programming

iterative approaches to find approximations for optimal value functions

1. policy evaluation: calculate value function with current policy (Bellman equation as update rule)
2. policy improvement: adjusting policy to act greedy (pick actions with maximum values) with respect to value function of current policy

putting both components together:

- policy iteration: $\pi_0 \xrightarrow{\text{E}} v_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} v_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \dots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} v_*$
- value iteration: truncated policy evaluation using Bellman optimality equation as update rule




from Sutton

Bootstrapping and Sampling

bootstrapping: update estimates of state values based on estimates of values of successor states

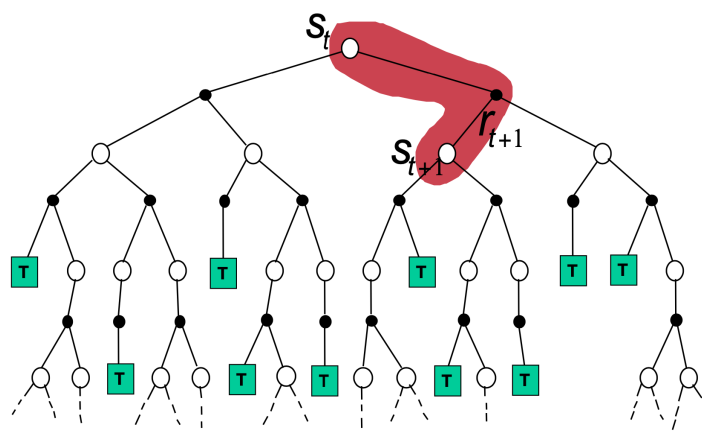
sampling: experience of sample sequences (no need for complete knowledge of environment)

Dynamic Programming



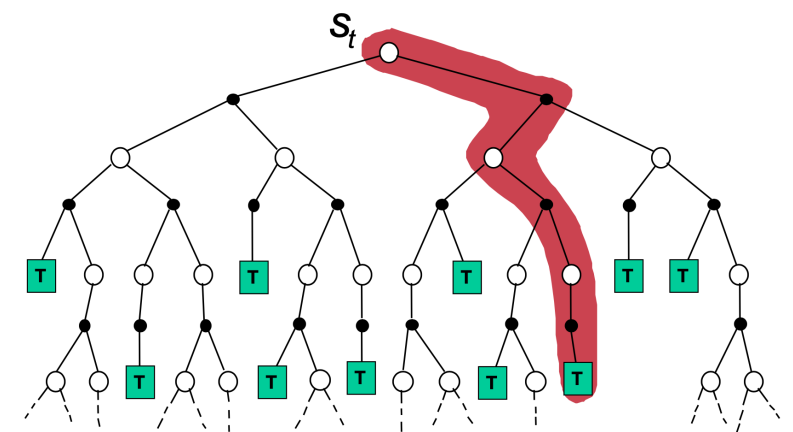
- bootstrapping
- no sampling \rightarrow model-based (transition probabilities needed)

Temporal Difference (TD) Learning



- bootstrapping
- sampling \rightarrow model-free

Monte Carlo (MC)



- no bootstrapping
- sampling \rightarrow model-free

from Sutton

Sampling Update Rule

$$NewEstimate \leftarrow OldEstimate + StepSize [Target - OldEstimate]$$

$$MC: \quad v(s_t) \leftarrow v(s_t) + \alpha [\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} - v(s_t)]$$

$$TD: \quad v(s_t) \leftarrow v(s_t) + \alpha [r_{t+1} + \gamma v(s_{t+1}) - v(s_t)]$$


bootstrapping



On-Policy TD Control: SARSA

generalized policy iteration:

estimating action-value function Q for current behavior policy π



change policy toward greediness with respect to Q (exploration for example via epsilon-greedy policy)

update:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$




Off-Policy TD Control: Q-Learning

estimate action-value function Q directly approximating optimal action-value function (independent of policy being followed \rightarrow potentially off-policy)

policy just determines which state-action pairs are visited and updated


update:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

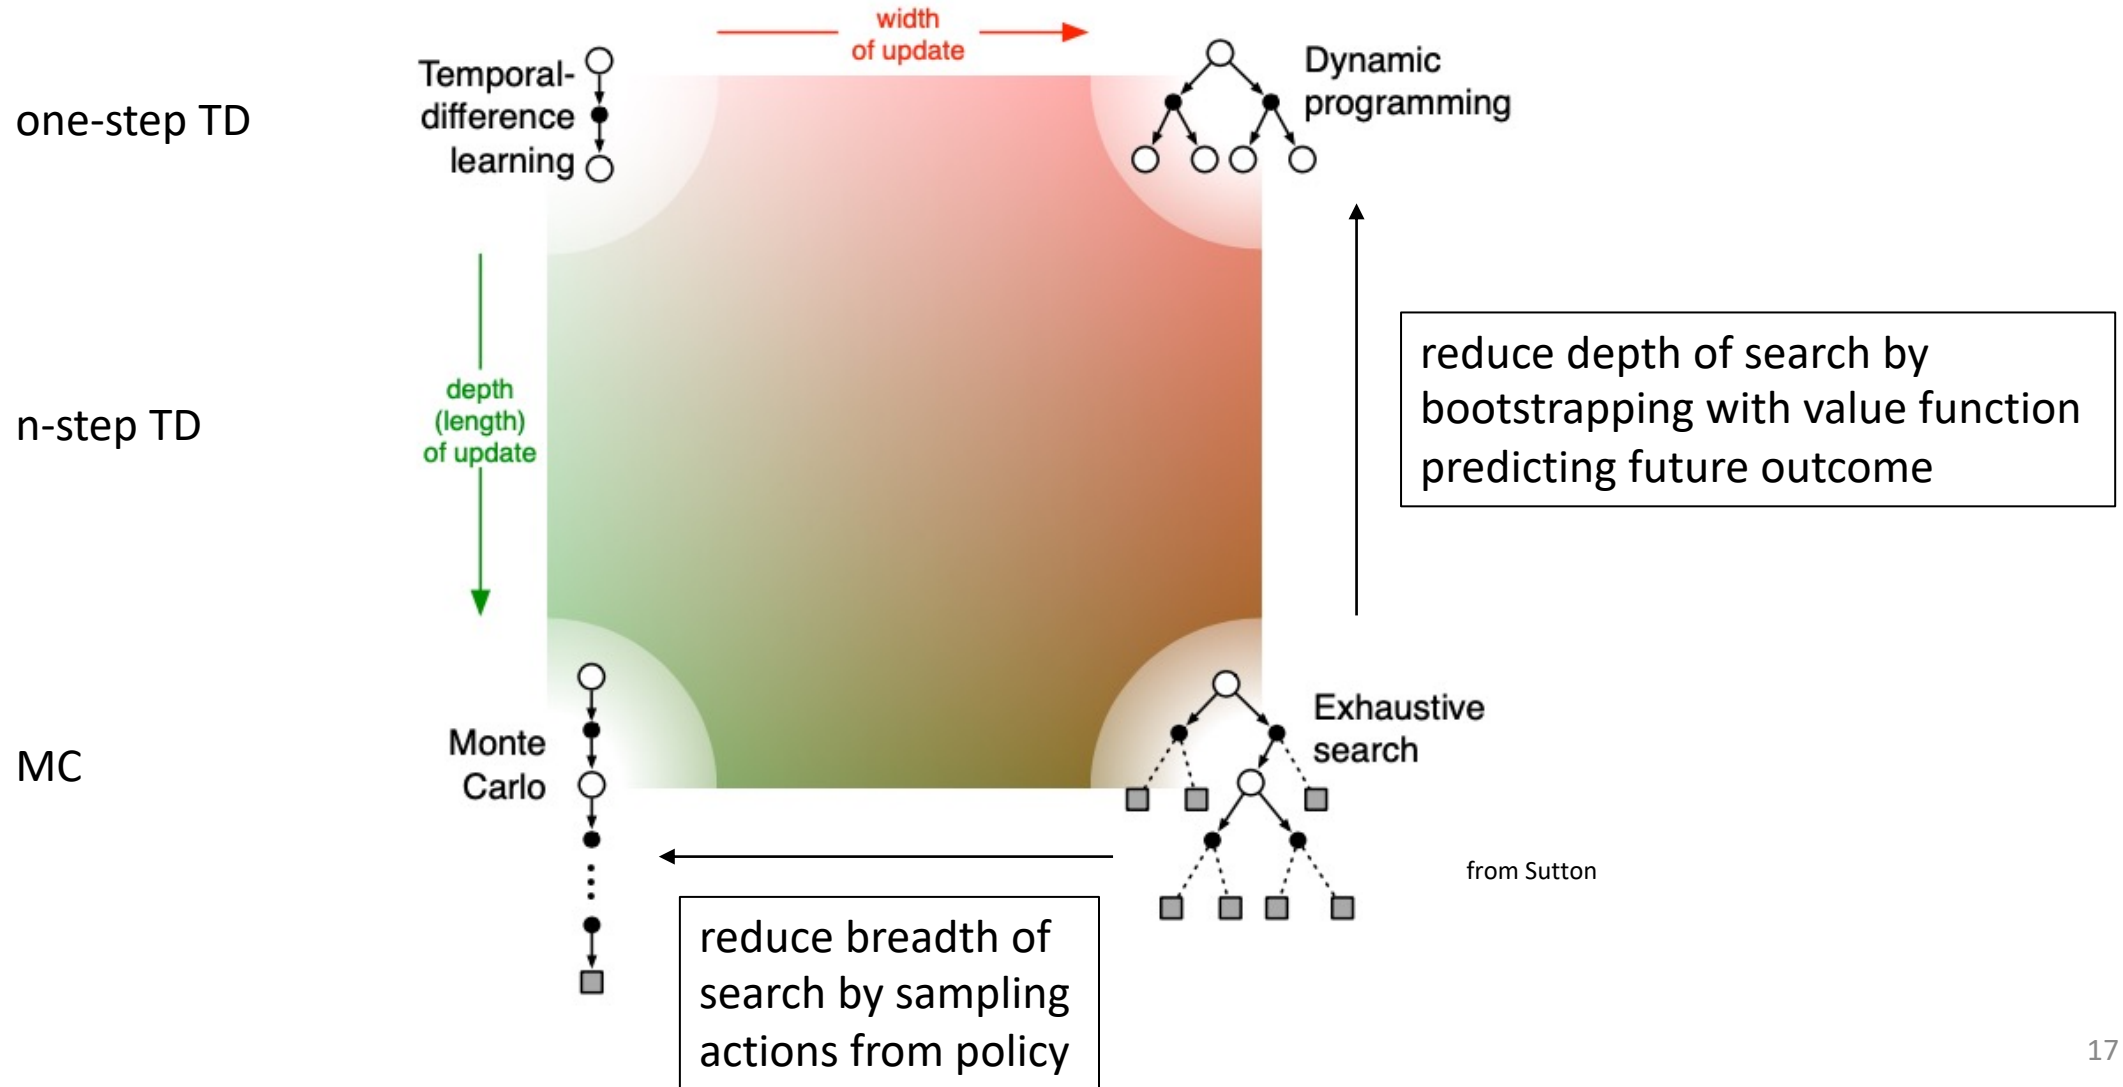


expected Sarsa:

$$\begin{aligned} Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \mathbb{E}_\pi [Q(S_{t+1}, A_{t+1}) \mid S_{t+1}] - Q(S_t, A_t) \right] \\ &\leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right] \end{aligned}$$



Summary: Update Characteristics



Deep Reinforcement Learning

Limitation of Tabular Methods

tabular methods simply memorize observed data

problem with tabular solution methods in practice: large state/action spaces → curse of dimensionality

need for generalization: supervised learning to the rescue

- non-linear function approximation
- nowadays often deep learning methods → deep reinforcement learning

Approximate Solution Methods

- value-function as parametrized functional form of state s with weight vector \mathbf{w} (instead of table)
- \mathbf{w} contains parameters for different features describing s (e.g., connection weights in neural network)

objective function
(mean squared value error):

$$\overline{\text{VE}}(\mathbf{w}) \doteq \sum_{s \in \mathcal{S}} \mu(s) \left[v_{\pi}(s) - \hat{v}(s, \mathbf{w}) \right]^2$$

μ : state distribution
(e.g., fraction of
time spent in s)

stochastic gradient descent:

$$\begin{aligned} \mathbf{w}_{t+1} &\doteq \mathbf{w}_t - \frac{1}{2} \alpha \nabla \left[v_{\pi}(S_t) - \hat{v}(S_t, \mathbf{w}_t) \right]^2 \\ &= \mathbf{w}_t + \alpha \left[v_{\pi}(S_t) - \hat{v}(S_t, \mathbf{w}_t) \right] \nabla \hat{v}(S_t, \mathbf{w}_t) \end{aligned}$$

with

$$\nabla f(\mathbf{w}) \doteq \left(\frac{\partial f(\mathbf{w})}{\partial w_1}, \frac{\partial f(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial f(\mathbf{w})}{\partial w_d} \right)^{\top}$$

Deep Q-Network (DQN)

deep neural network to approximate Q-function (Q-value as output for any state-action pair)

Mnih et al. (Google DeepMind): *Human-level control through deep reinforcement learning*

separate target network (weights only periodically updated with Q-network weights)

→ reducing correlations of Q-network with target

experience replay: apply Q-learning updates on samples/minibatches of experience drawn at random from pool of stored samples (agent's experiences at each time-step)

→ removing correlations in observation sequence (make it i.i.d.)

Side Note: ...

... i.i.d. as fundamental assumption of ML

... i.i.d. \rightarrow causality

Famous Example of Deep RL: AlphaGo

Monte Carlo tree search (heuristic search algorithm) for move (action) selection

guided by deep convolutional neural networks for both value function and policy estimation

→ improving search efficiency

reduce depth of search tree by evaluating positions with **value function**
(predicting outcome from given position → **bootstrapping**)

reduce breath of search tree by **sampling** actions using **policy network**
(probability distribution over possible moves in given position)

Side Note: Model-Predictive Control

... beam-search-based planning conceptually an instance of model-predictive control

Direct Policy Search

Policy Gradient Methods

learning of parametrized policy (without value functions):


$$\pi(a|s, \boldsymbol{\theta}) = \Pr\{A_t = a \mid S_t = s, \boldsymbol{\theta}_t = \boldsymbol{\theta}\}$$

parameters: e.g., neural network weights

maximizing objective $J(\boldsymbol{\theta})$ (expected cumulative rewards)

update rule of REINFORCE method:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \alpha \cdot \underbrace{\nabla(\log \pi(A_t | S_t, \boldsymbol{\theta}))}_{\nabla J(\boldsymbol{\theta})} \cdot \underbrace{G_t}_{\text{“weighting” with return}}$$



policy gradients $\nabla \pi$:
e.g., neural network gradients

Actor-Critic Methods

hybrid between policy-based and value-based methods (to reduce variance)

value function as critic of policy (instead of return):

$$\theta_{t+1} = \theta_t + \alpha \cdot \nabla(\log \pi(A_t | S_t, \theta)) \cdot Q(S_t, A_t)$$

independent parametrizations for π and Q (e.g., two separate neural networks)

advantage actor-critic: $Q(S_t, A_t) \rightarrow A(S_t, A_t) = Q(S_t, A_t) - V(S_t)$

$\underbrace{Q(S_t, A_t) - V(S_t)}$
can be approximated
by TD error

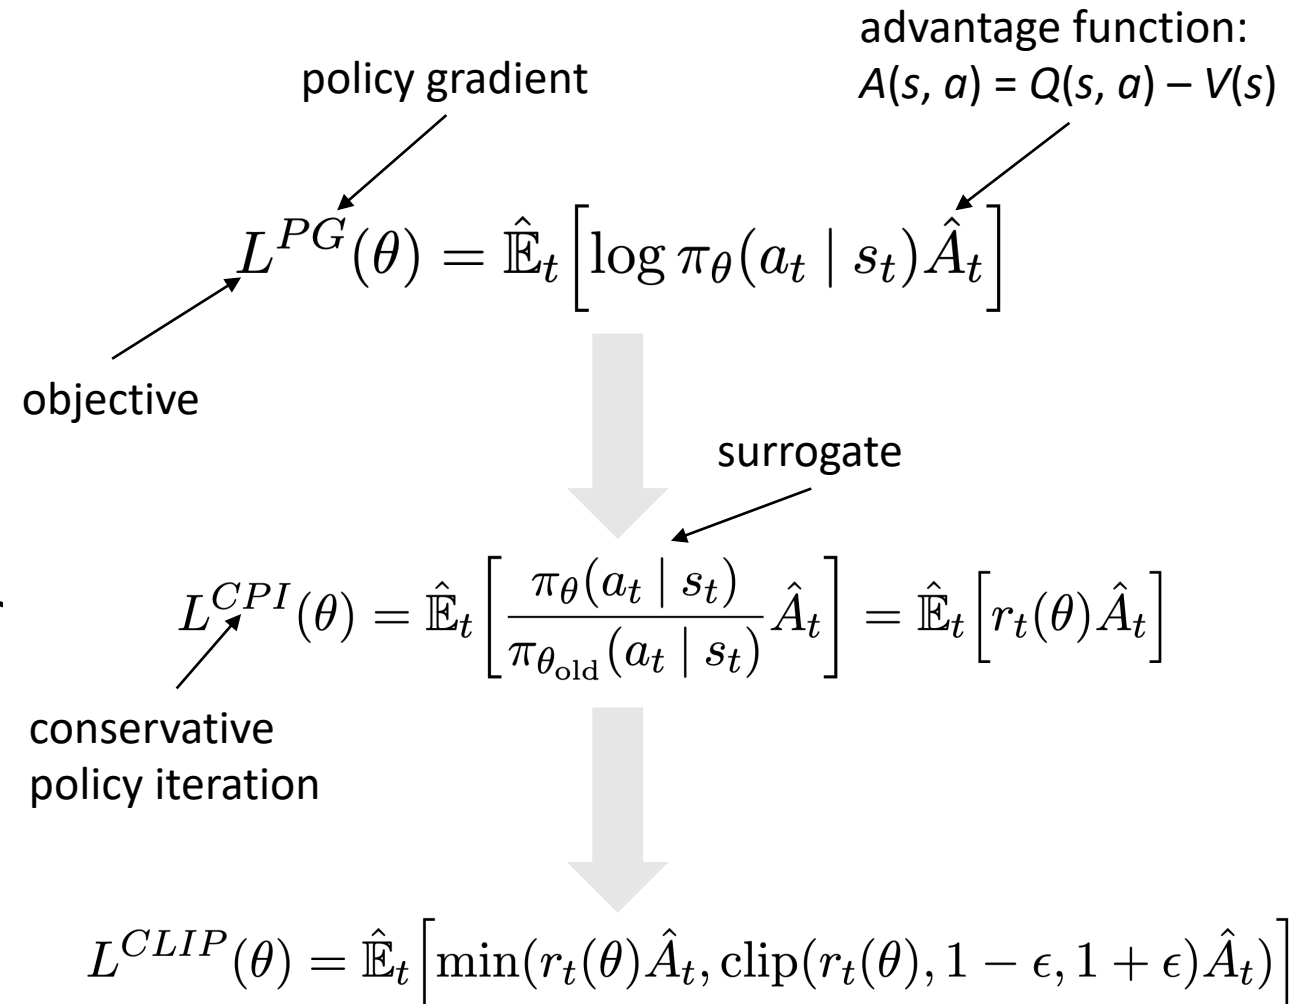
Proximal Policy Optimization (PPO)

state-of-the-art policy gradient method

advantage actor-critic method with clipped surrogate objective function

- surrogate objective from trust region policy optimization → better efficiency
- clipping: limiting policy update at each training step → improved stability of actor

trust-region methods: first choose size of trust region, then direction
line-search methods: first choose direction, then step size

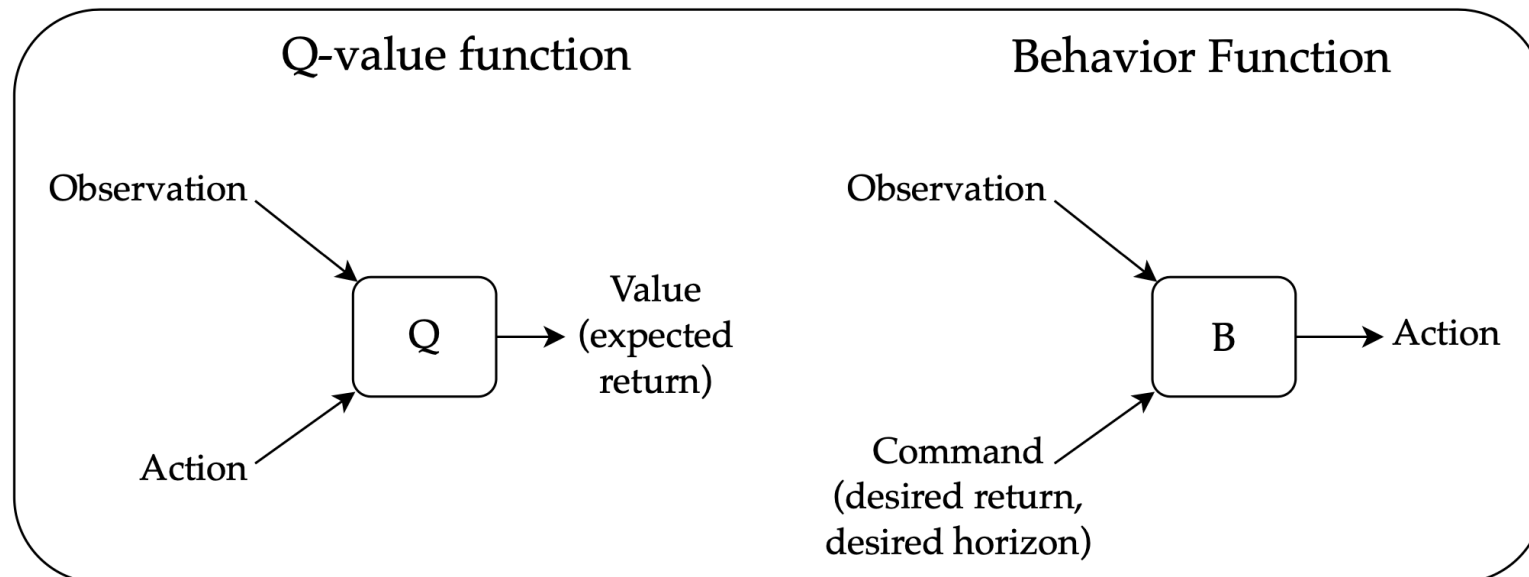


Upside-Down RL

combination of off-policy bootstrapping (e.g., Q-learning) with high-dimensional function approximation leads to non-stationary targets (deadly triad)

most popular technique to overcome this: target networks (a copy of an agent's value function is frozen and stored periodically to provide stationary learning targets for temporal-difference learning)

upside –down RL as alternative



Generative Trajectory Modeling

transformer (sequence model) trained on fixed, limited experience consisting of trajectory rollouts of arbitrary policies (offline RL)

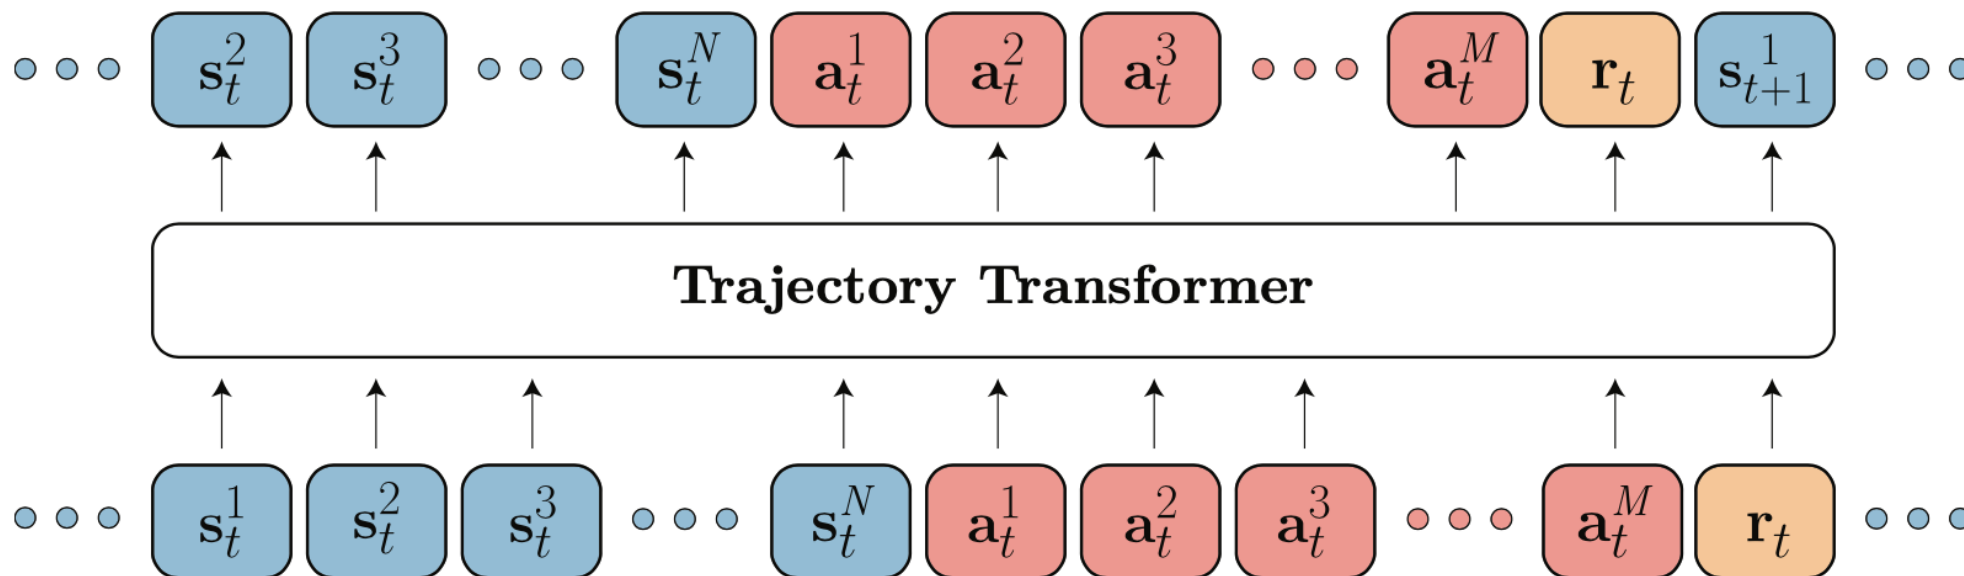
→ no need for bootstrapping

perform credit assignment directly via self-attention: implicitly forming state-return associations via similarity of query and key vectors (maximizing the dot product)


decoder architecture to autoregressively model trajectories

- Trajectory Transformer: sequence model for joint distribution of states, actions, and rewards
- Decision Transformer: conditional sequence model, conditioning on desired return (reward), past states, and actions to generate future actions

planning mirrors sampling procedure used to generate sequences from language model: selecting desired return tokens, acting as prompt for generation



Literature



papers:

- ...

Automation

...one of most impactful goals of AI

...computer vision, NLP

next step:

automated decision-making/control (e.g., autonomous driving)

...but also ... [nuclear fusion plasma stabilization](#)

...control, robotics