

Machine Learning – Overview

June 2023

Main Areas of Artificial Intelligence

- **computer vision**
(spatial structures, state-of-the-art: Convolutional Neural Networks)
- **natural language processing**
(sequential structures, state-of-the-art: transformers)
- **automated decision making, robotics**
(reinforcement learning)



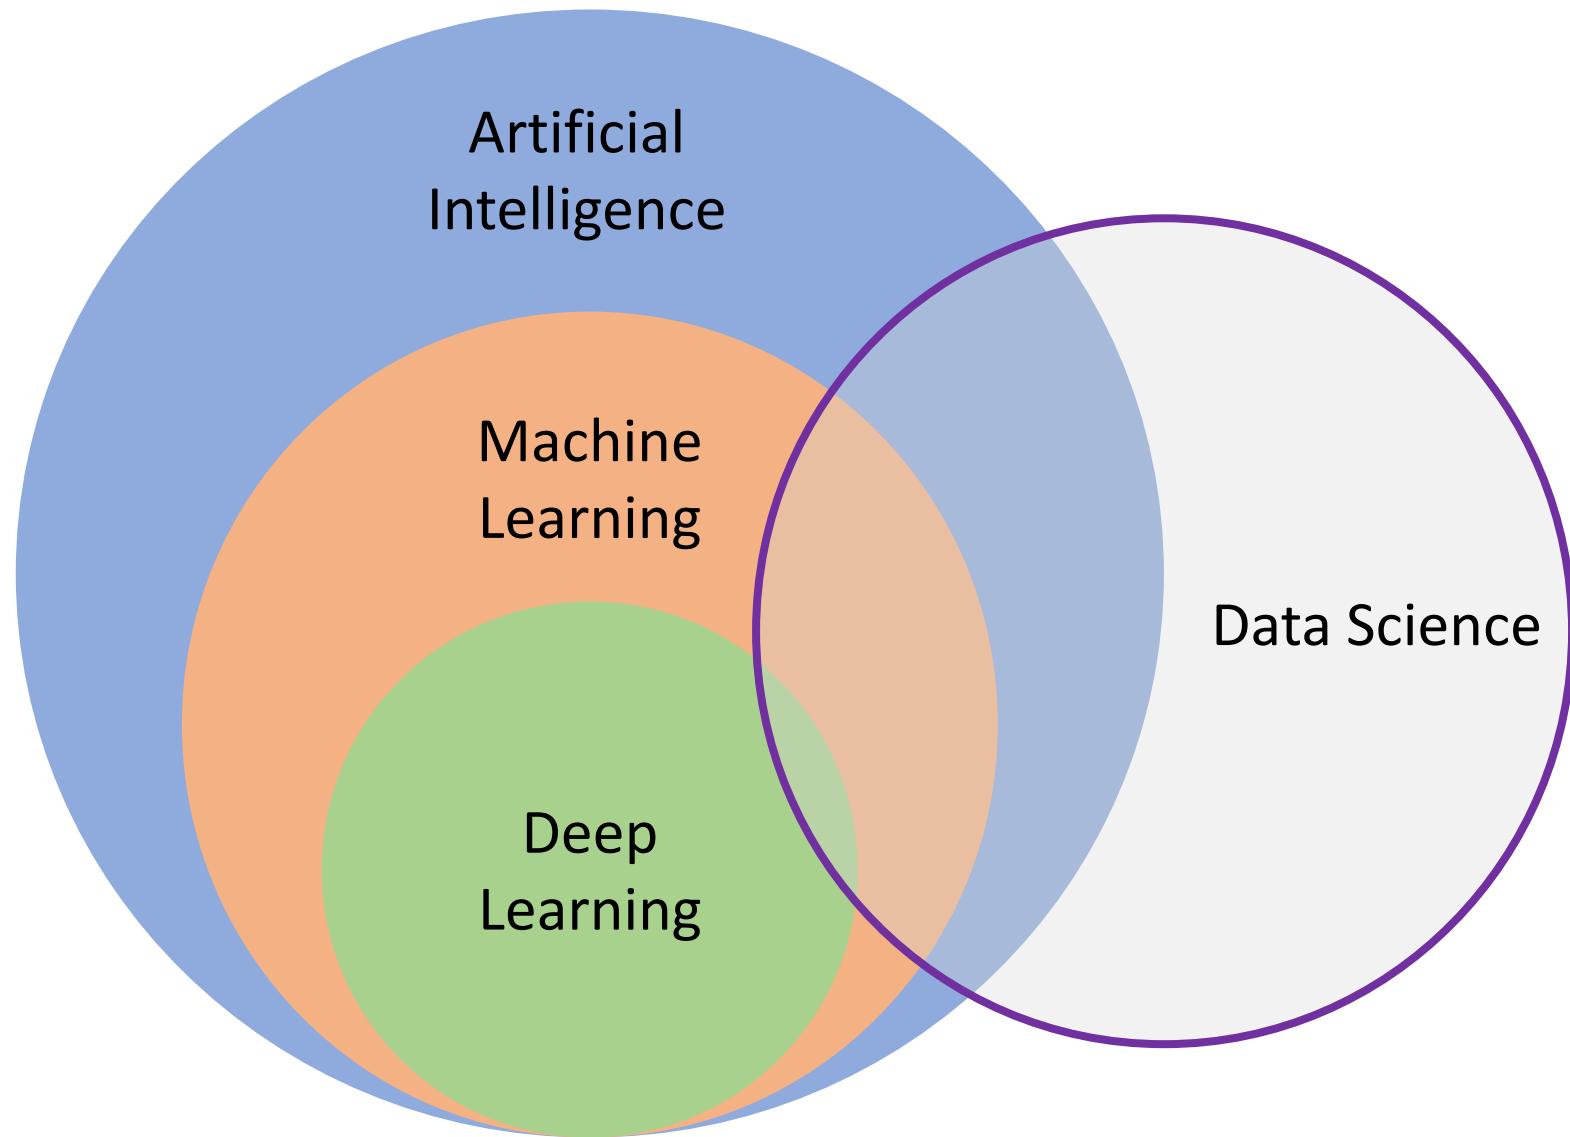
from wikipedia

All of these are enabled by one key ingredient:

- *learning from experience (Machine Learning)*
- knowledge representation, automated reasoning

agency:
perception – thought – action

Buzz Words ...



Deep Learning:
special kind of ML
algorithms using (deep)
neural networks

Data Science:
extract knowledge from
data (by means of ML,
among other things)

Traditional Algorithms and GOFAI

traditional algorithms:

explicit (handcrafted) instructions for each situation



from wikipedia

symbolic AI (aka GOFAI):

use knowledge by means of symbols (as representations), logic, search
(e.g., expert systems like Deep Blue)

Public perception is changing over time: A modern chess program, nowadays disparaged as brute computing, would have been considered intelligent in the 50s.

ML: Learning from Experience/Data

mainly exploiting statistical dependencies with the aim of **generalization** to new (e.g., future) data

training (usually offline optimization):

ML algorithm + data = explicit algorithm (to be used at inference time)

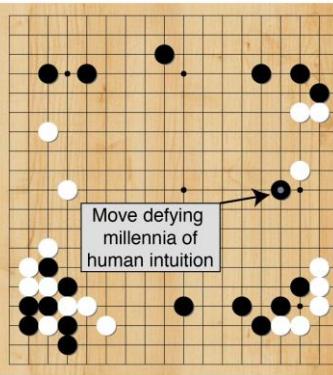
→ reduction of complexity and much better generalizability compared to handcrafted algorithms

analogy: Humans do not hit the ground running (storage capacity of DNA limited), but have learning capabilities.

Hybrid Approaches for ML and Symbolic AI

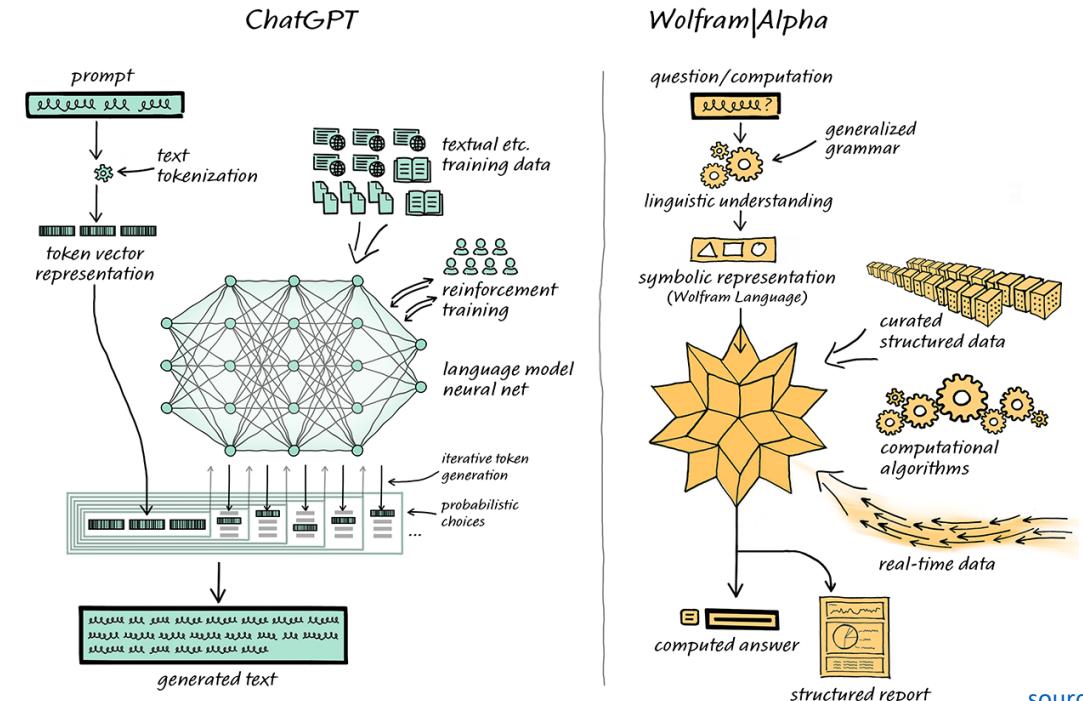


vs.
AlphaGo



famous move 37
from Max Tegmark

Deep Reinforcement Learning
and Monte Carlo Tree Search



feature engineering for ML models also
kind of symbolic knowledge representation

tool usage:
[LangChain](#)

Supercharging the Scientific Method

use ML and data to replace or enhance explicit methods relying on detailed domain knowledge ([Software 2.0](#))

- overcome our evolutionary limitations in math with clever learning algorithms and collecting data
- immediate impact on many aspects of industry, business, and science, formulated as narrow tasks with strictly defined inputs (aka weak AI)

more imminent than (still philosophical) long-term quest for human-level AI (aka strong AI, AGI), i.e., general-purpose intelligence
(although recent language models show multi-purpose capabilities)

When to apply ML?

complexity

- decisions under uncertainty, many influencing factors
- e.g., demand forecasting, DNA sequencing
- difficult for humans, direct model inexpressible

automation

- e.g., face and speech recognition, autonomous driving
- goal to reach human-level performance

... and of course you need data to learn from

and more recently:
generative tasks

- rather than predictive (or discriminative) ones
- e.g., image generation, conversational AI, new proteins or materials

Learning Paradigms

Supervised Learning

learning by teacher → usually rather narrow tasks

Target Quantity

- **known in training:** labeled samples or observations from past
- to be for unknown cases (e.g., future values)

Features

- input information that is
- correlated to target quantity
 - known at prediction time



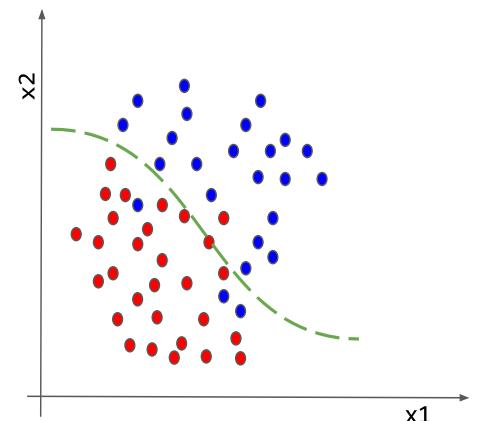
Example: Spam Filtering

Classify emails as spam or no spam

use accordingly **labeled emails as training set**

use information like **occurrence of specific words or email length** as **features**

features x_1 and x_2
spam, no spam



Reinforcement Learning

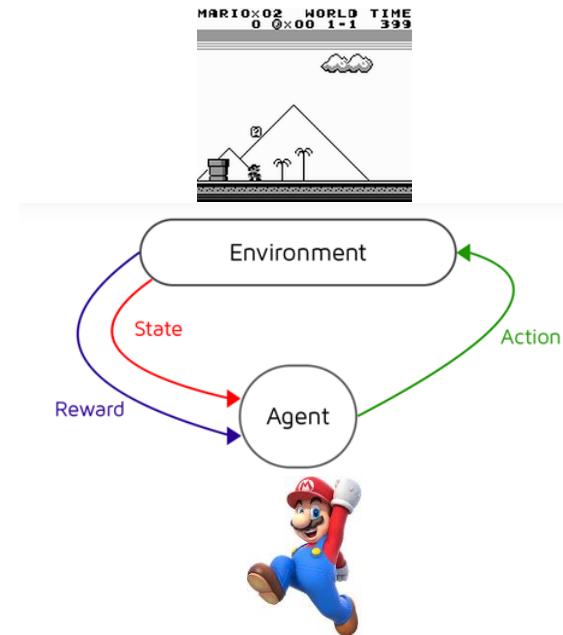
learning by trial-and-error (exploration and exploitation)

- goal-based approach → more generic than supervised learning (but sparse reward signals)
- receiving feedback from the environment, no supervision
- formalization of sequential decision making (delayed rewards)

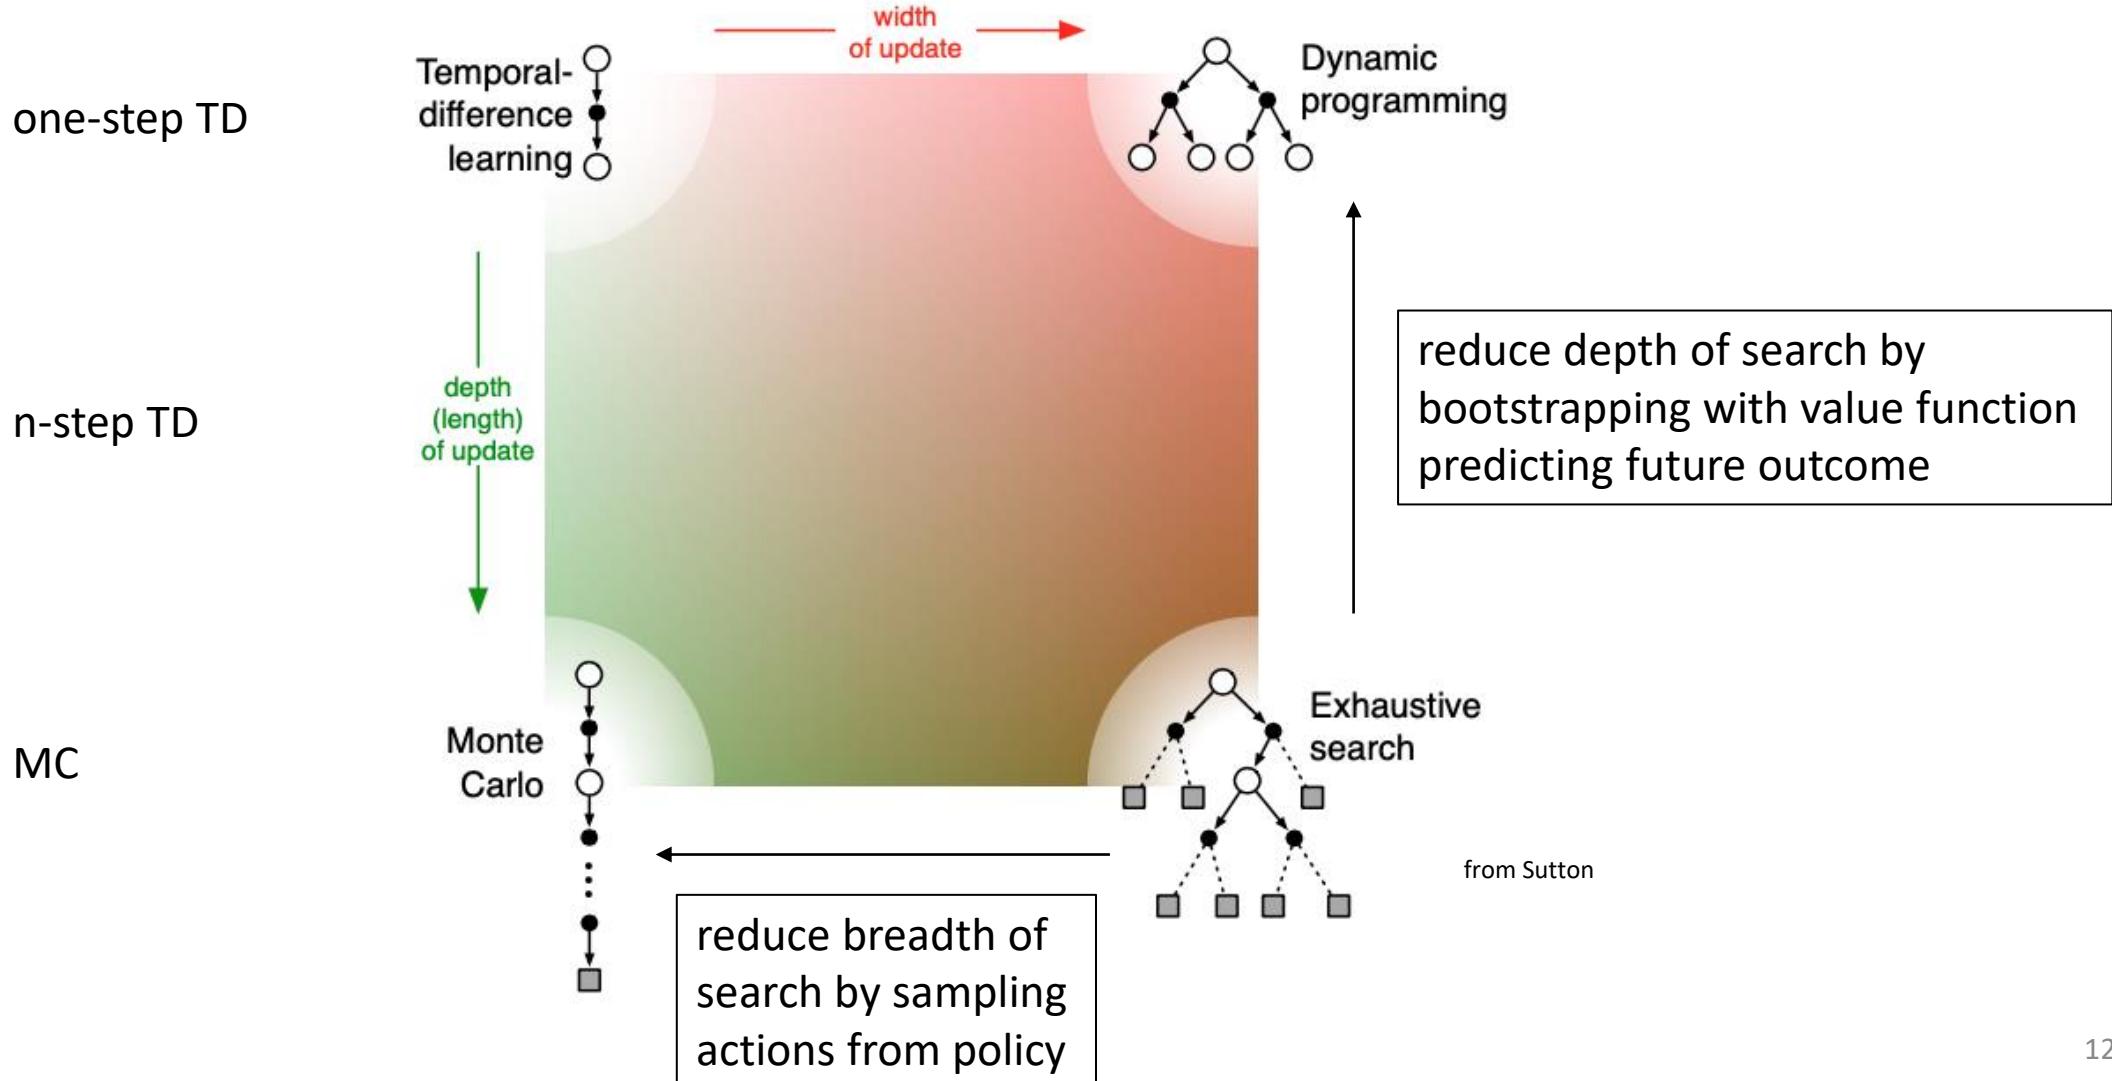
corresponds to search for best action policy to reach a given goal (e.g., win a game)

using learning from examples (data) to guide the search

RL usually more difficult (e.g., non-differentiable as a whole) than supervised learning (which can be seen as “generalized optimization”, often of proxy metric)



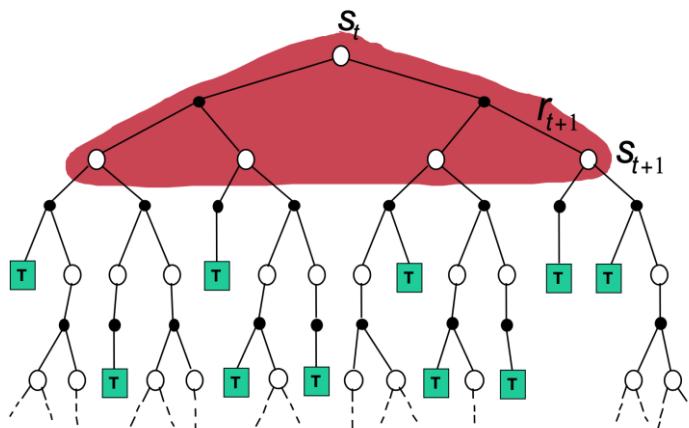
Reduction of Search Space



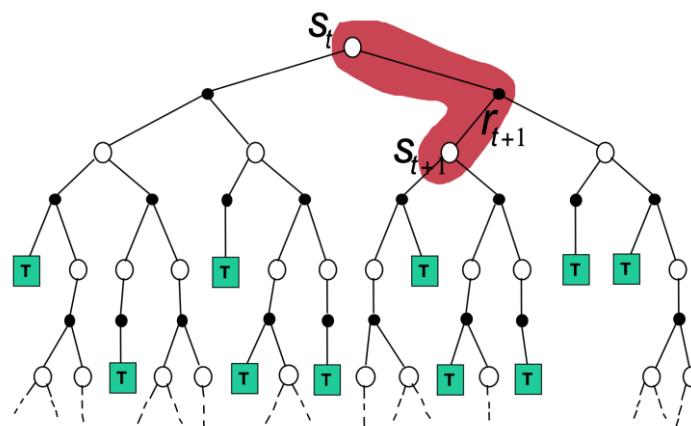
Main Concepts of Value-Based RL Methods

bootstrapping: update estimates of state values based on estimates of values of successor states
sampling: experience of sample sequences (no need for complete knowledge of environment)

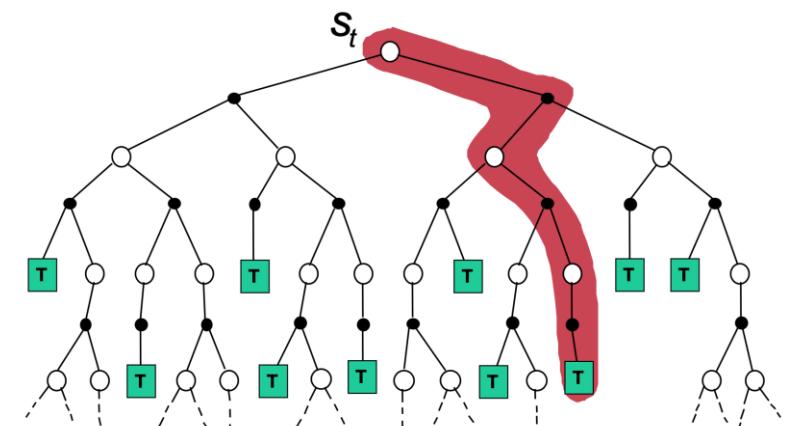
Dynamic Programming



Temporal Difference (TD) Learning



Monte Carlo (MC)



from Sutton

- bootstrapping
- no sampling → model-based
(transition probabilities needed)

- bootstrapping
- sampling → model-free

- no bootstrapping
- sampling → model-free

Unsupervised Learning

learning by observation

no target information → kind of “vague” pattern recognition (but plenty of data)

self-supervised:

- input-output mapping like supervised learning
- but generating labels itself from input information
- learning of semantic feature representations
- e.g., word2vec, BERT, GPT

ML needs lots of training data

How Much Information is the Machine Given during Learning?

- ▶ “Pure” Reinforcement Learning (**cherry**)
 - ▶ The machine predicts a scalar reward given once in a while.
- ▶ **A few bits for some samples**

- ▶ Supervised Learning (**icing**)
 - ▶ The machine predicts a category or a few numbers for each input
 - ▶ Predicting human-supplied data
 - ▶ **10→10,000 bits per sample**

- ▶ Self-Supervised Learning (**cake génoise**)
 - ▶ The machine predicts any part of its input for any observed part.
 - ▶ Predicts future frames in videos
 - ▶ **Millions of bits per sample**



© 2019 IEEE International Solid-State Circuits Conference

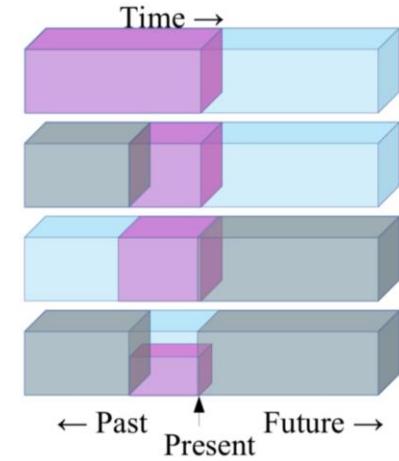
1.1: Deep Learning Hardware: Past, Present, & Future

Y. LeCun

59

Self-Supervised Learning

- ▶ Predict any part of the input from any other part.
- ▶ Predict the **future** from the **past**.
- ▶ Predict the **future** from the **recent past**.
- ▶ Predict the **past** from the **present**.
- ▶ Predict the **top** from the **bottom**.
- ▶ Predict the **occluded** from the **visible**
- ▶ **Pretend there is a part of the input you don't know and predict that.**



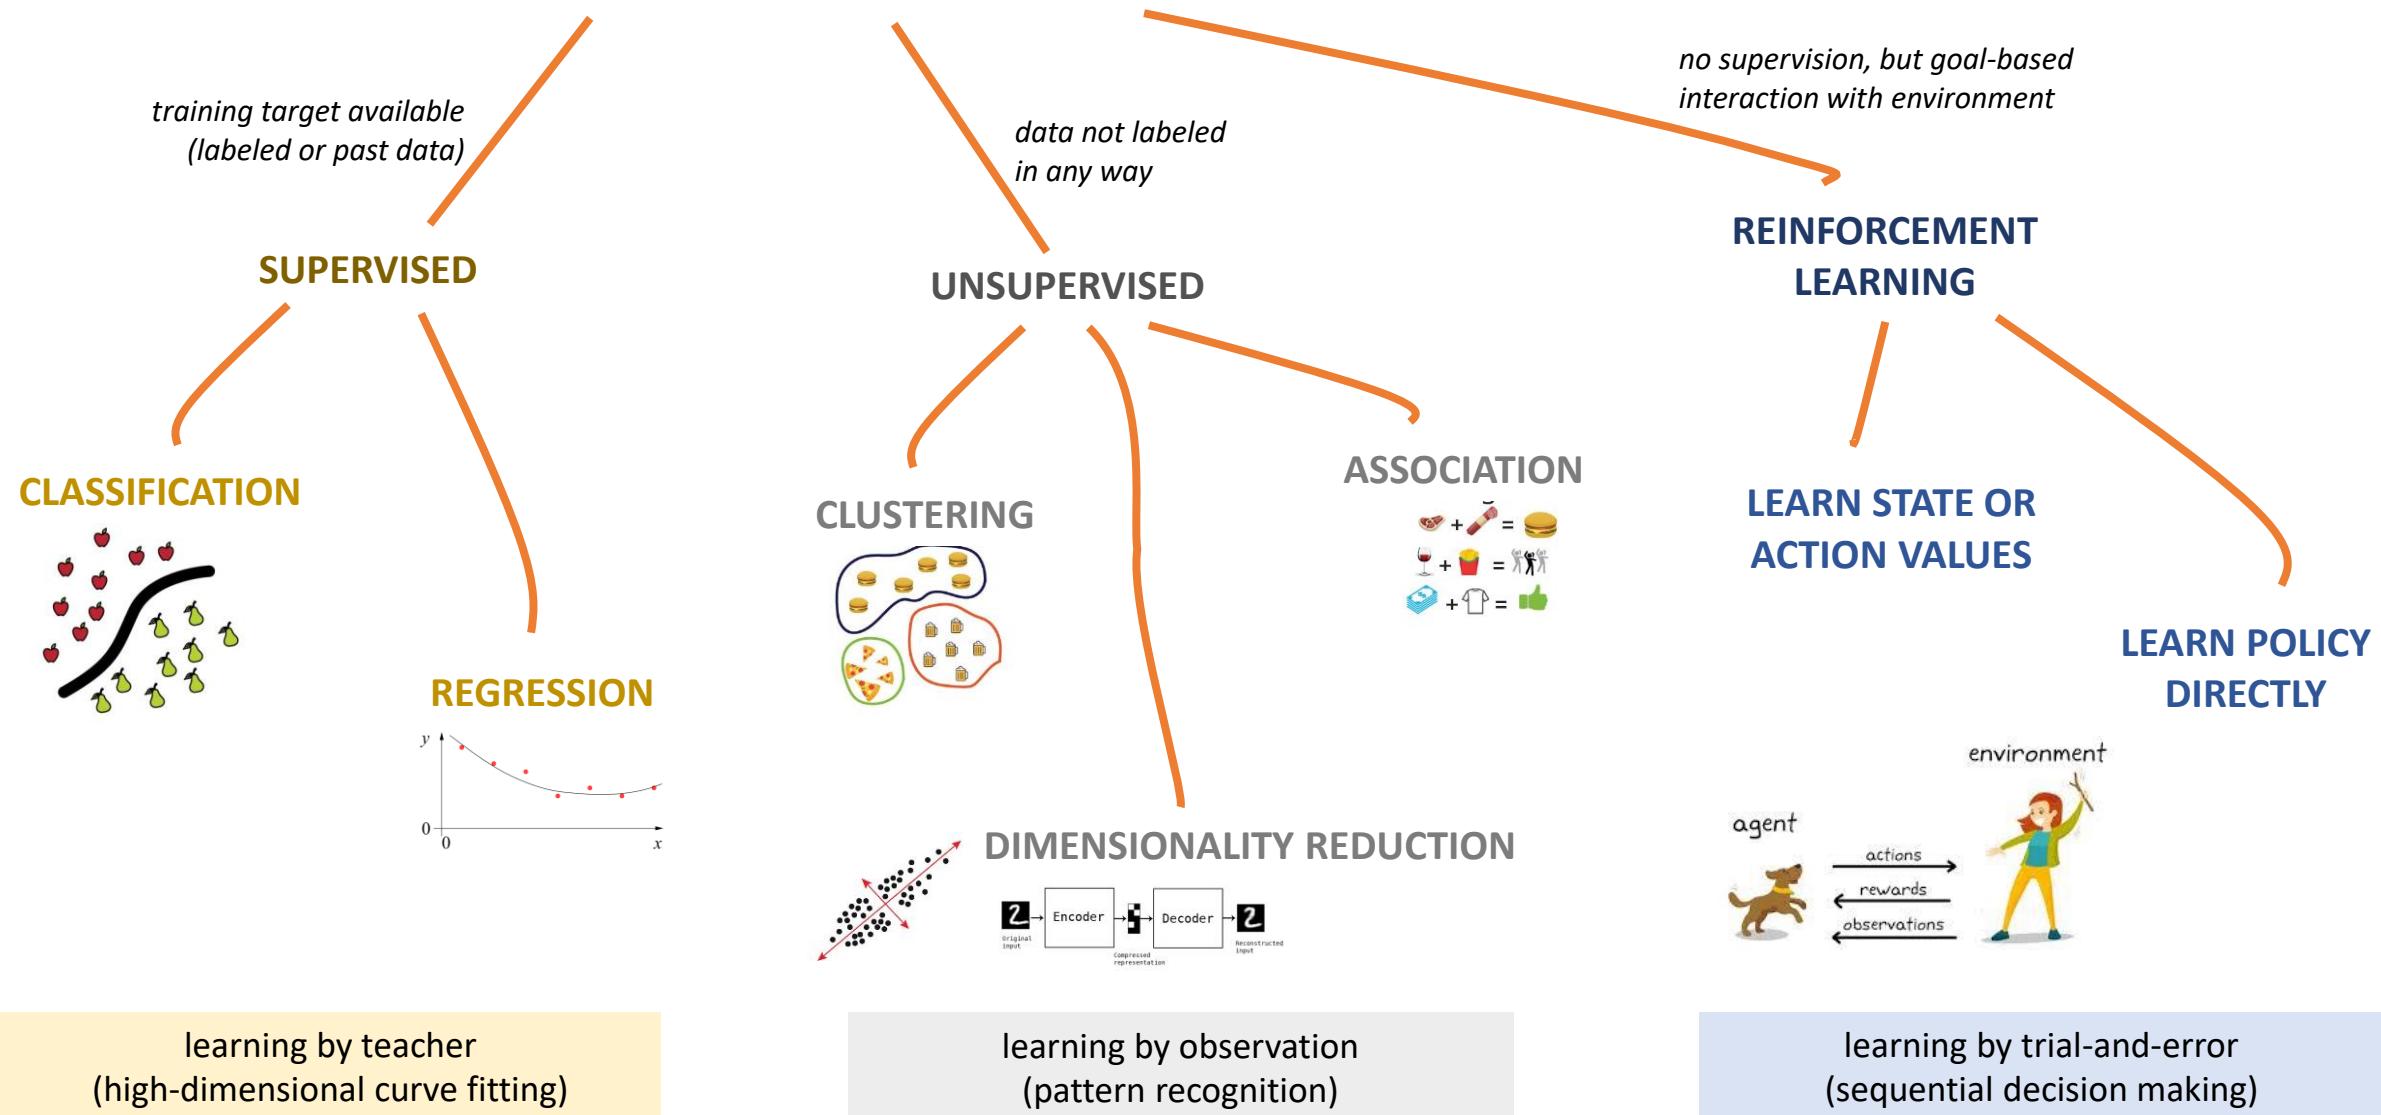
© 2019 IEEE International Solid-State Circuits Conference

1.1: Deep Learning Hardware: Past, Present, & Future

Y. LeCun

58

MACHINE LEARNING



Fitting and Generalization



... ML ...



Supervised Learning Scenario

map inputs to output: $y = f(\mathbf{x})$ (estimated: $\hat{f}(\mathbf{x})$)

random variables Y and $\mathbf{X} = (X_1, X_2, \dots, X_p)$ ← usually many dimensions

fit train data set of (y_i, \mathbf{x}_i) pairs

(i.i.d. assumption: random samples from underlying data-generating process)

then apply learned statistical dependencies to test data set

classification:

categorical target: $y = 0$ or $y = 1$ (e.g., image of cat or not), predict probabilities

regression:

real-valued target: $Y \in [0, \infty)$ (e.g., demand forecasting) or $Y \in (-\infty, \infty)$

Generalization

core of ML:

empirical risk minimization (training error) as proxy for minimizing unknown population risk (test error, aka generalization error or out-of-sample error)

generalization gap: difference between test and training error

interpolation: to unencountered samples from training environment

extrapolation: to testing conditions differing from training environment

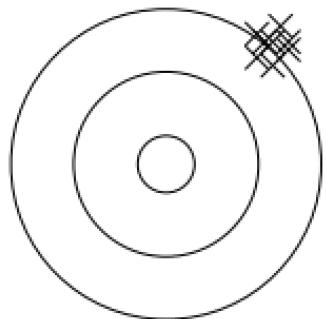
curse of dimensionality: “*learning in high dimensions always amounts to extrapolation*”

but reality is friendly: most high-dimensional data sets reside on lower-dimensional manifolds (manifold hypothesis) → enabling effectiveness of ML

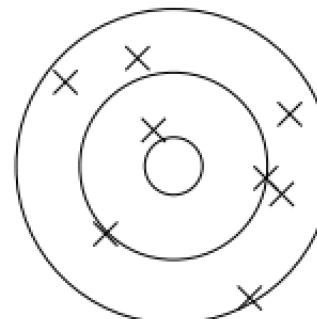
→ need for appropriate **inductive bias** (aka learning bias): set of assumptions of a learning algorithm to predict outputs of inputs not encountered during training (“data in disguise”)

Bias, Variance, Irreducible Error

think of fitting ML algorithms as repeatable processes with different (i.i.d.) data sets



bias:
due to too simplistic model
(same for all training data sets)
“underfitting”



variance:
due to sensitivity to specifics (noise)
of different training data sets
“overfitting”

irreducible error (aka Bayes error):

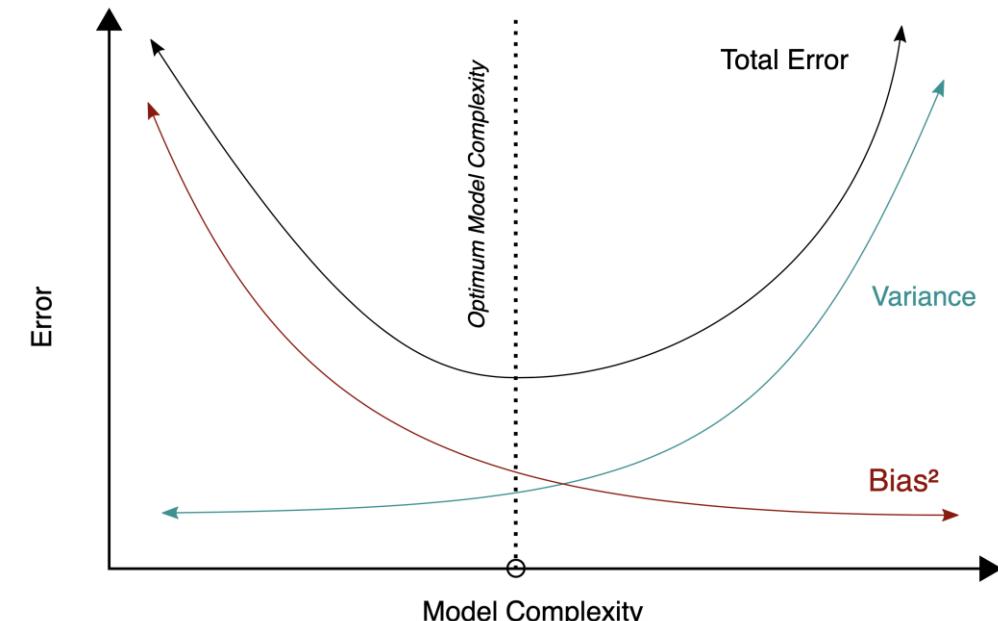
inherent randomness (target generated from random variable following probability distribution)

→ limiting accuracy of ideal model

different potential reasons for inherent randomness (noise): complexity, missing information, ...

Bias-Variance Tradeoff

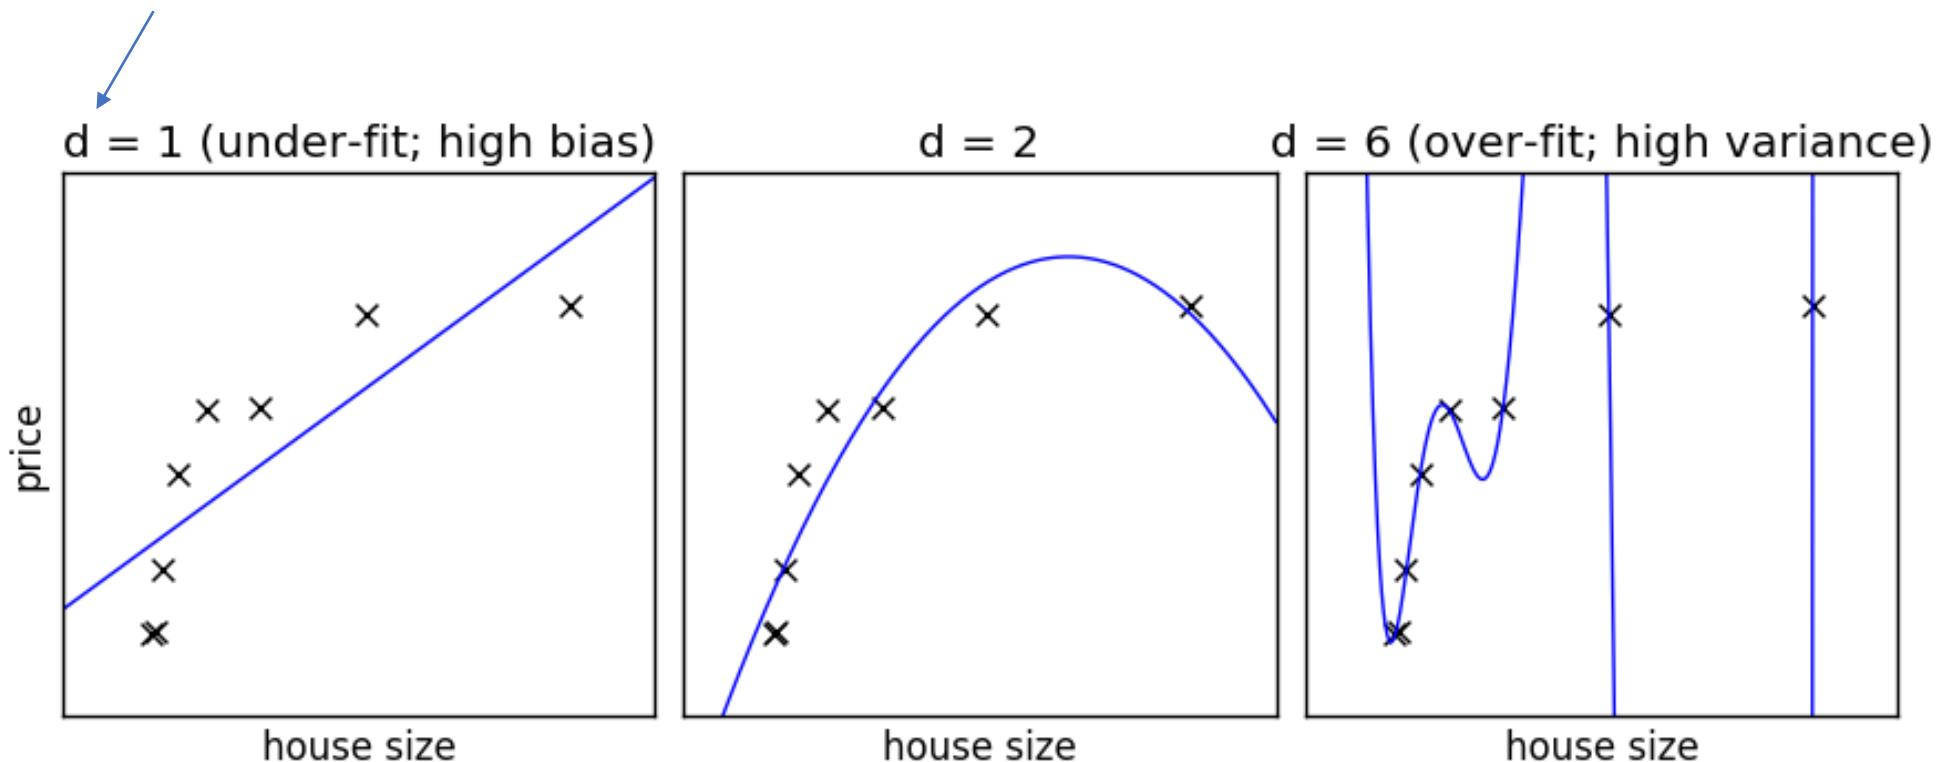
- fundamental concept in classical statistical learning theory
- models of higher complexity have lower bias but higher variance (given the same number of training examples)
- generalization error follows U-shaped curve:
overfitting once model complexity (number of parameters) passes certain threshold
- overfitting: variance term dominating test error
→ increasing model complexity increases test error



from wikipedia

Example: Non-Linear Function Approximation

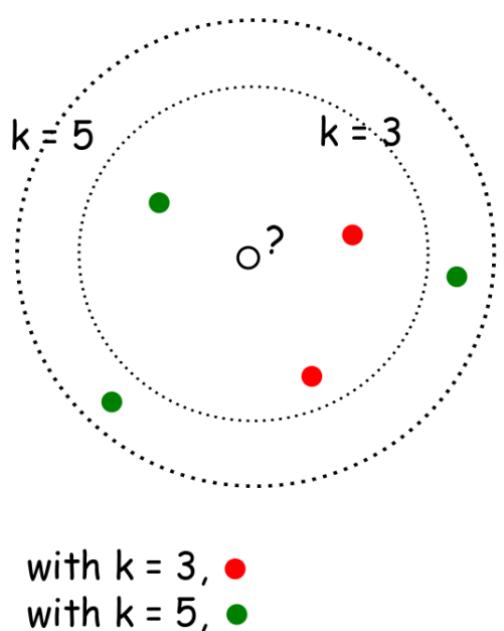
degree of fitted polynomial



from scikit-learn documentation

Example: k-Nearest Neighbors

- local method, instance-based learning
- non-parametric
- distance defined by metric on x (e.g., Euclidean)



regression:

$$\hat{f}(x_0) = \frac{1}{k} \sum_{j=1}^k y_j \quad \text{with } j \text{ running over } k \text{ nearest neighbors of } x_0$$

- low k : low bias but high variance
- high k : low variance but high bias

$$bias = f(x) - \frac{1}{k} \sum_{j=1}^k y_j$$

$$var = \frac{\sigma^2}{k}$$

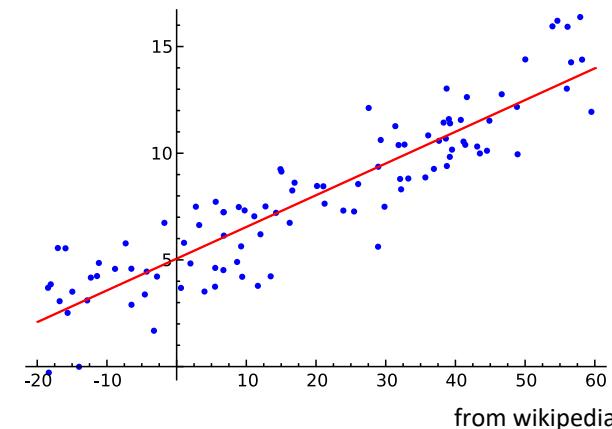
Linear Regression

fit:

$$y_i = \hat{\alpha} + \sum_{j=1}^p \hat{\beta}_j x_{ij} + \varepsilon_i$$

(model)

error term (noise): reflects assumed data distribution (here: Gaussian with same variance σ^2 for all samples)



parameters to be estimated:

- $\hat{\alpha}, \hat{\beta}$

$$\rightarrow \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

(approximating assumed true α, β, σ)

predict:

$$\hat{y}_i = E[Y|X = \mathbf{x}_i] = \hat{f}(\mathbf{x}_i)$$

- conditional mean for squared loss of least squares method
- predict arbitrary quantile by means of quantile loss

$$p(y|\mathbf{x}_i) = \mathcal{N}(y; \hat{y}_i, \hat{\sigma}^2)$$

Gaussian mean variance
(reflected by ε_i in fit)

General Recipe of Statistical Learning

statistical learning algorithm by combining:

- **model** (e.g., linear function, Gaussian distribution)
- **objective function** (e.g., squared residuals)
- **optimization algorithm** (e.g., gradient descent)
- **regularization** (e.g., convolutions)

Loss Function

loss function L : expressing deviation between prediction and target

$$L(y_i, \hat{f}(x_i); \hat{\theta})$$

with $\hat{\theta}$ corresponding to parameters of model $\hat{f}(x)$

e.g., $\hat{\alpha}, \hat{\beta}$ in linear regression

e.g., squared residuals (for regression problems):

$$L(y_i, \hat{f}(x_i); \hat{\theta}) = (y_i - \hat{f}(x_i; \hat{\theta}))^2$$

Cost Function

averaging losses over (empirical) training data set:

$$J(\hat{\theta}) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(x_i); \hat{\theta})$$

cost function to be minimized according to model parameters $\hat{\theta}$
→ objective function

Cost Minimization

minimize training costs $J(\hat{\theta})$ according to model parameters $\hat{\theta}$:

$$\nabla_{\hat{\theta}} J(\hat{\theta}) = 0$$

e.g., for mean squared error (aka least squares method):

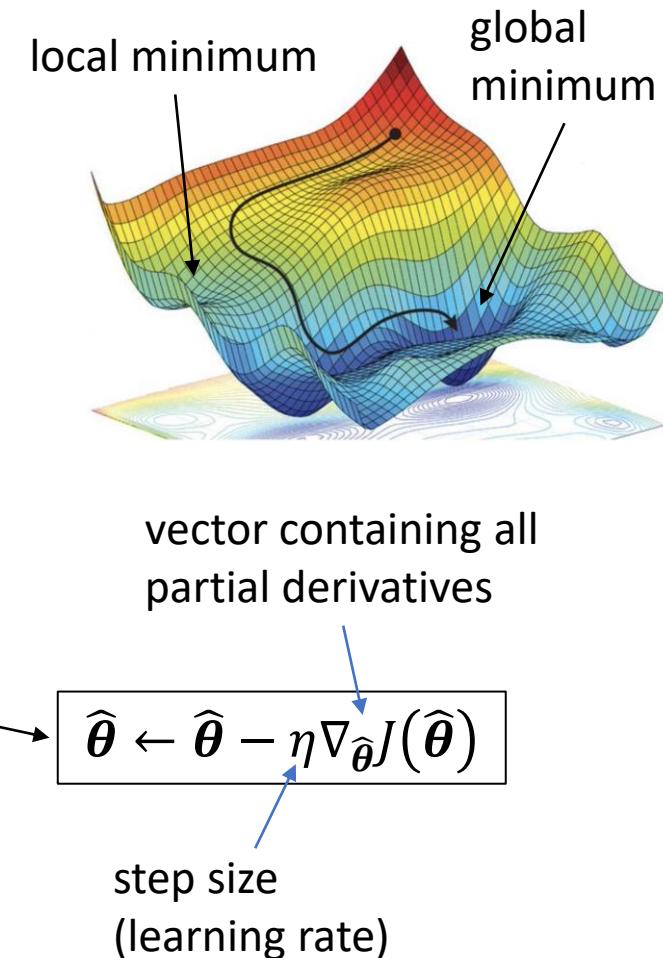
$$\nabla_{\hat{\theta}} \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i; \hat{\theta}))^2 = 0$$

analytical solution for linear regression: ordinary least squares

in general: iterative, numerical optimization (e.g., **gradient descent**)

maximum likelihood estimation (minimization of D_{KL} between probability distributions of true data-generating process and model:
make the model distribution match the empirical distribution):

special objective function, estimate mode of assumed model distribution

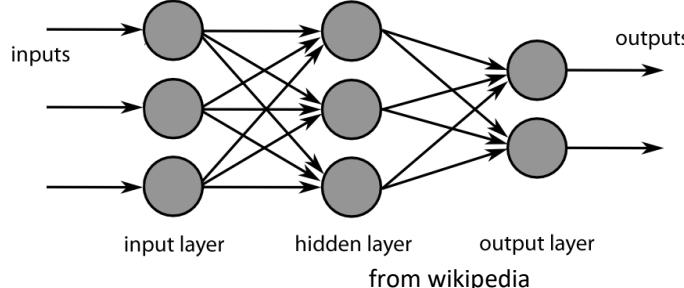
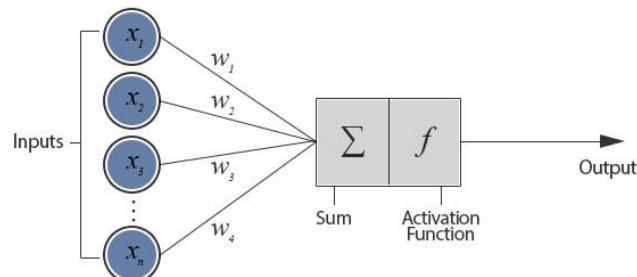


Algorithmic Families

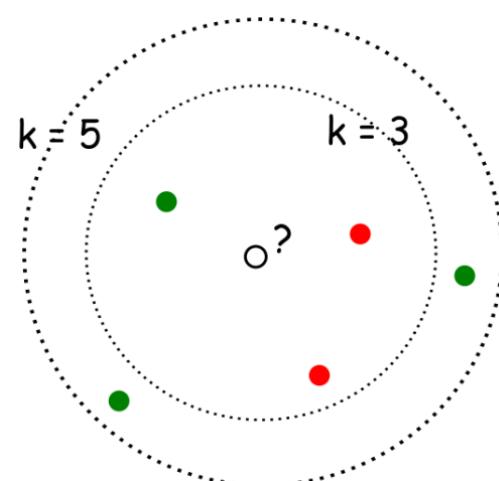
linear (parametric) models

- linear regression
- Generalized Linear Models
- Generalized Additive Models

neural networks: non-linear just by means of activation functions



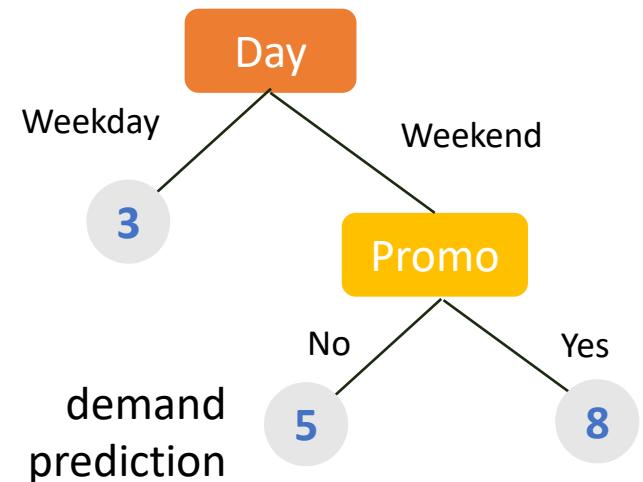
nearest neighbors (local methods, instance-based learning) – non-parametric models



with $k = 3$, ●
with $k = 5$, ●

kernel/support-vector machines: linear model (maximum-margin hyperplane) with kernel trick

decision trees



often used in ensemble methods

- bagging: random forests
- boosting: gradient boosting

Most ML algorithms can be described by the general recipe of combining models, costs, optimization, and regularization methods, including non-linear models like neural networks (backpropagation), support-vector machines (hinge loss in soft-margin SVM), or decision trees (impurity functions), and even unsupervised learning (e.g., maximum variance axes in PCA)

Most powerful ML algorithms are compound, with rather simple (often linear) building blocks.

To generalize well, one needs to find a method with an appropriate inductive bias for the task at hand (e.g., regularization method like convolutional layers or objective function).

Deep Learning

Recap: Goal of ML

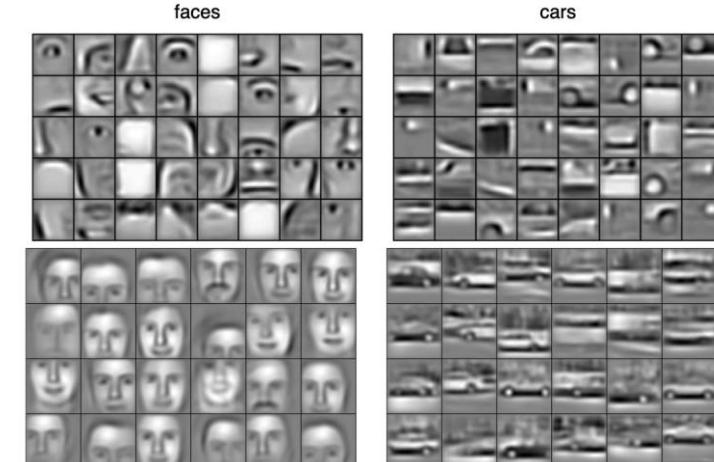
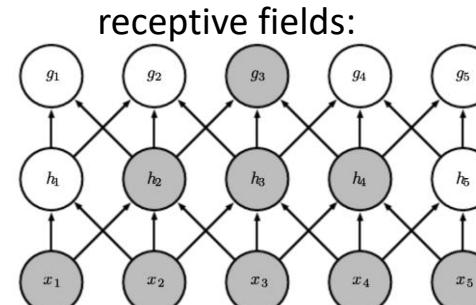
generalization from optimization on training data set (approximation of true data generating probability distribution by empirical risk minimization)

- fitting: complex function approximation
- for generalization: learning of good abstraction/representation of data/concepts

→ deep learning methods (neural networks with many layers) optimal candidates

e.g., convolutional neural networks (CNN) can learn hierarchical representation by means of many convolutional and pooling layers

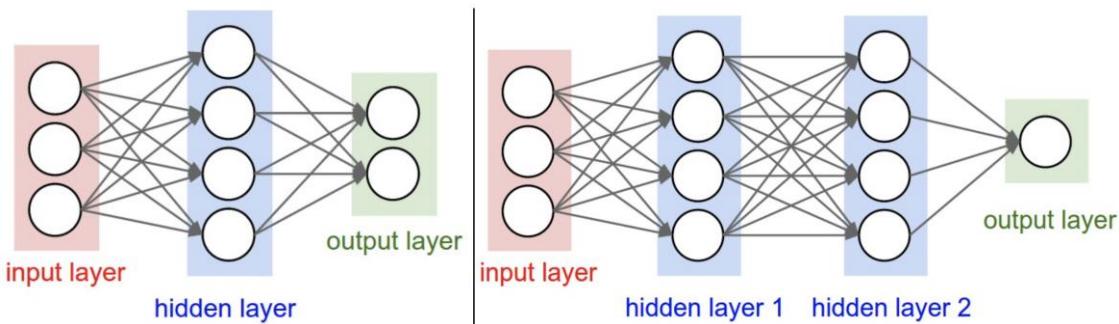
the deeper the better:
accuracy, hierarchical representation



[source](#)

Multi-Layer Perceptron (MLP)

fully-connected feed-forward network with at least one hidden layer



toward deep learning:
add more nodes and hidden layers ...

classification:

- logistic regression in hidden nodes
- cross-entropy loss: $L_i(y_i, \hat{f}(\mathbf{x}_i); \hat{\mathbf{w}}) = -\sum_{k=1}^K y_{ik} \log \hat{f}_k(\mathbf{x}_i; \hat{\mathbf{w}})$
- several output nodes k for multi-classification
- softmax output function: $g_k(\mathbf{t}_i) = \frac{e^{t_{ik}}}{\sum_{l=1}^K e^{t_{il}}}$

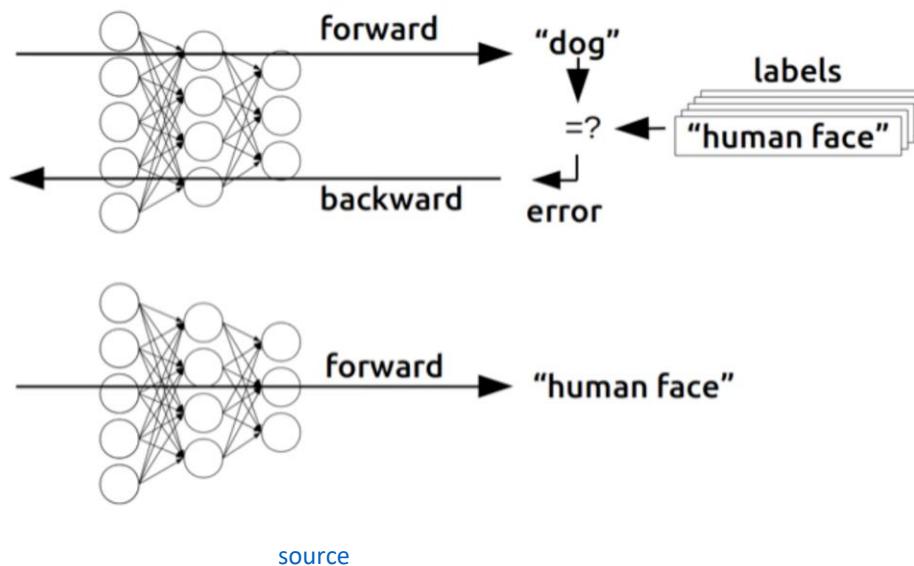
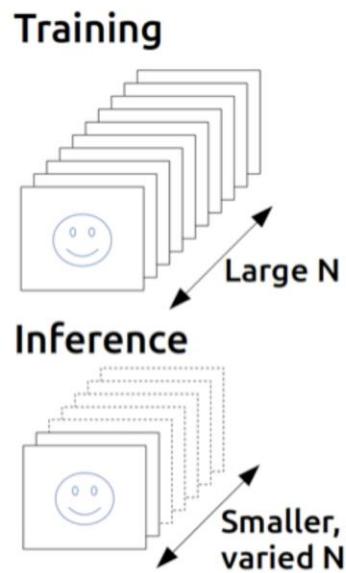
regression:

- squared error loss
- identity output function
- usually just one output node

Learning Mechanism: Back-Propagation

back-propagation of errors (gradients of cost function according to weights) through layers via chain rule of calculus (avoiding redundant calculations of intermediate terms)

each node exchanges information only with directly connected nodes → enables efficient, parallel computation



- forward pass: current weights fixed, predictions computed
- backward pass: errors computed from predictions and back-propagated → weights then updated according to loss gradients (via gradient descent)

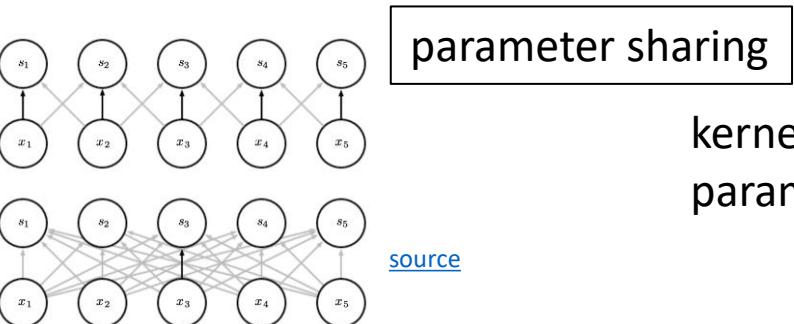
Convolutional Neural Networks

0	2	15	0	0	11	10	0	0	0	9	9	0	0
0	0	4	6	13	26	255	255	377	95	91	0	0	29
0	10	16	28	254	254	245	245	250	222	103	10	0	0
0	14	170	205	255	255	244	254	255	245	255	251	174	0
0	92	255	255	255	251	211	211	141	111	125	215	235	49
11	217	243	255	153	33	226	53	2	0	10	13	232	255
18	223	252	254	12	0	7	7	0	7	20	27	25	36
14	245	255	252	21	11	9	3	0	3	236	243	255	37
0	87	250	250	248	214	60	0	11	2	25	252	248	146
0	13	111	255	255	245	255	251	255	248	252	240	208	36
1	0	51	251	251	245	251	245	255	241	251	167	0	7
0	0	4	58	251	251	246	253	252	250	11	0	1	0
0	0	4	37	255	255	255	248	252	255	244	255	186	10
0	22	206	252	246	251	241	109	0	25	245	255	255	194
0	111	255	252	255	158	24	0	6	255	232	230	256	0
0	218	251	250	137	7	11	0	0	2	25	250	250	175
0	17	255	255	103	9	20	0	13	3	138	251	245	61
0	10	251	241	255	230	98	55	193	217	248	255	252	54
0	18	148	250	255	247	255	255	255	249	255	240	255	72
0	0	73	1	215	255	250	254	255	249	248	244	14	12
0	0	6	1	52	153	153	173	255	252	257	37	0	4
0	0	5	5	0	0	0	0	1	4	0	6	6	0

source

many images
(training examples),
potentially with
several channels

convolutional vs fully-connected layer:



kernel (learned
parameters)

parameter sharing

Convolution

Poolin

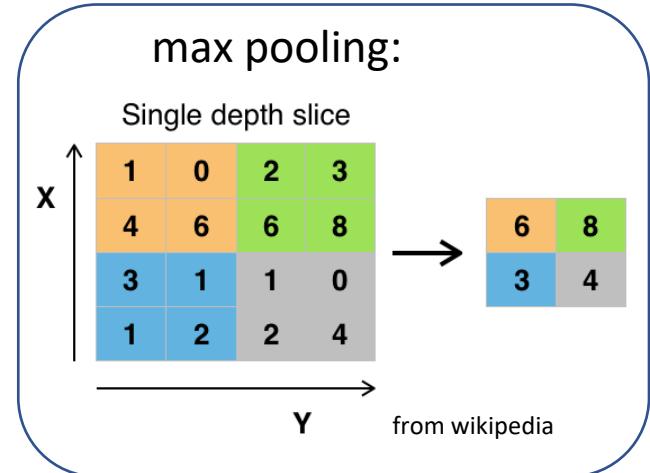
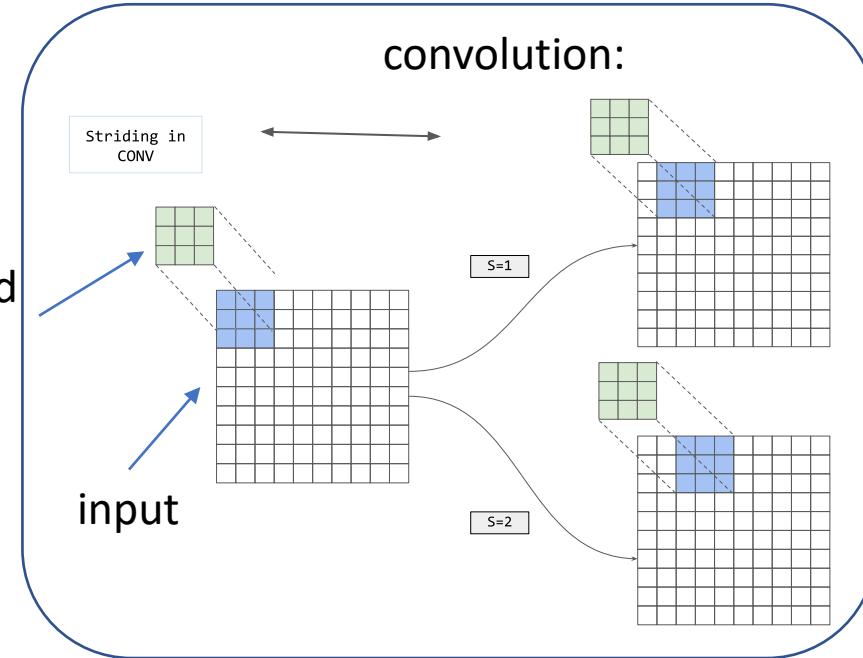
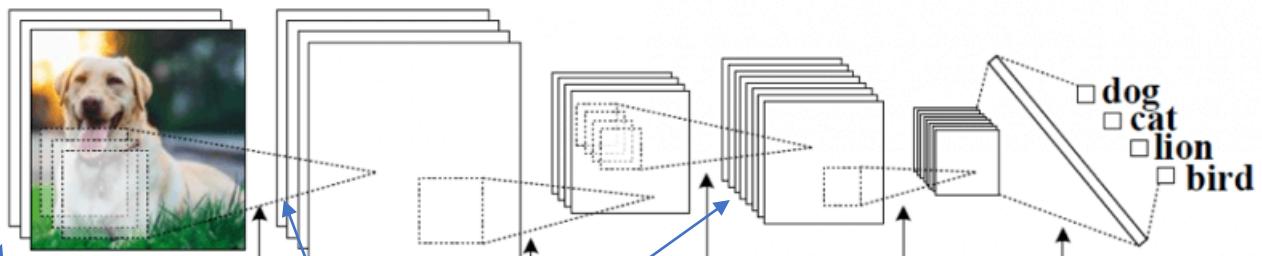
Involution Properties

~~Using Fully-connected~~

down-sampling
by convolutions
and pooling

several kernels
producing several
feature maps

flatten dimensions
for final classification
or regression



highly regularized form of feed-forward neural networks

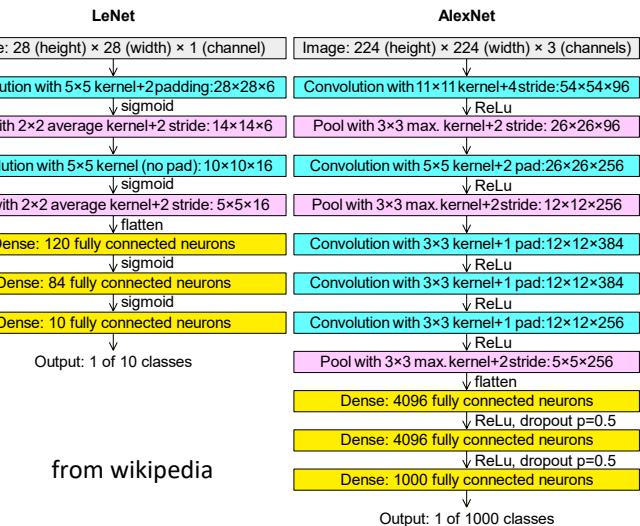
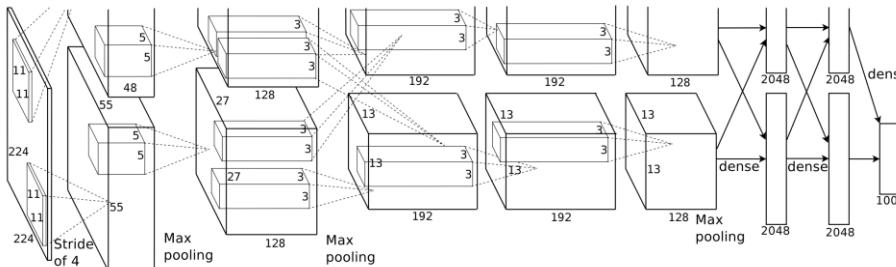
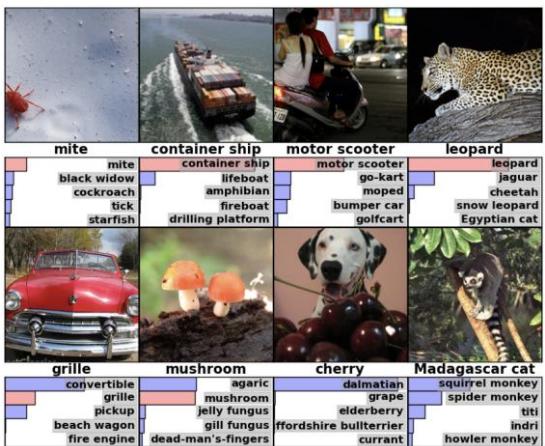
Rise of Deep Learning

a little bit oversimplified:

deep learning = lots of training data + parallel computation + smart algorithms

AlexNet: ImageNet + GPUs (allowing more layers) + ReLU, dropout

(pivotal moment for deep learning: ImageNet challenge 2012)



Training Subtleties of Deep Neural Networks

optimization and regularization difficult

- non-convex optimization problem (e.g., local vs global minima), easily overfitting
- many hyperparameters to tune

many methods to get it working in practice (despite partly patchy theoretical understanding)

optimization

- activation and loss functions
- weight initialization
- stochastic gradient descent
- adaptive learning rate
- batch normalization

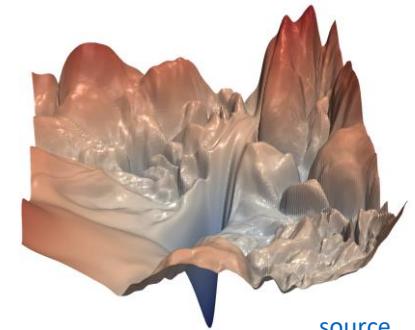
explicit regularization

- weight decay
- dropout
- data augmentation
- weight sharing

implicit regularization

- early stopping
- batch normalization
- stochastic gradient descent

typical loss surface:



[source](#)

Large Language Models (LLM)

natural language processing: dealing with sequential structures (e.g., text)
examples:

- machine translation (sequence-to-sequence model)
- sentiment classification
- chat bot (conversational AI)

context awareness via **embeddings** and (formerly) recurrent neural networks
(RNN) or (nowadays) **self-attention**

LLMs: **transformer** models with hundreds of billions of parameters

Embeddings

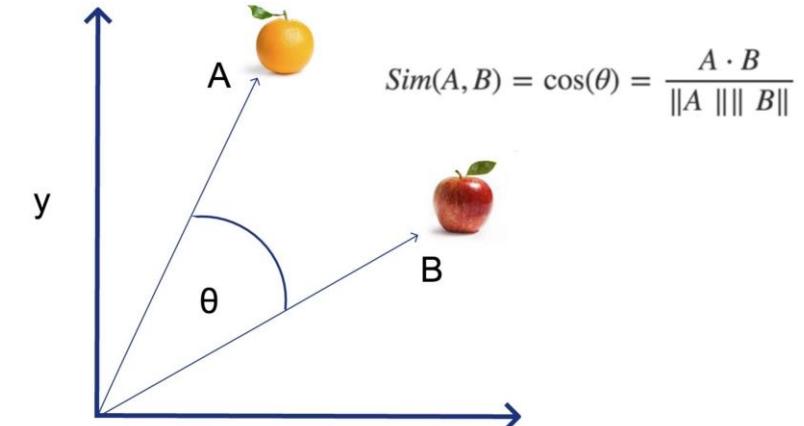
representation of entities by vectors

similarity between embeddings by, e.g., cosine
similarity → semantic similarity

most famous application: word embeddings
→ associations (natural language processing)

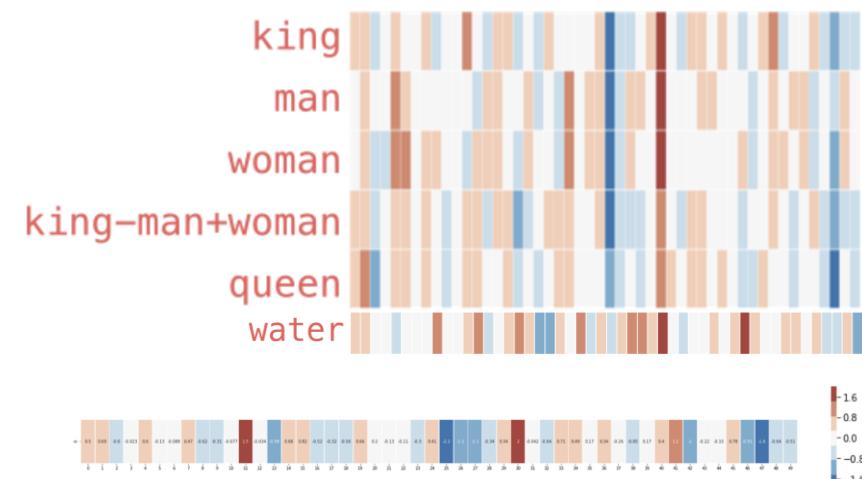
but general concept: embeddings of
(categorical) features (e.g., products in
recommendation engines)

learned via co-occurrence (e.g., [word2vec](#))



but also direction of difference
vectors interesting (analogies):

king – man + woman ≈ queen



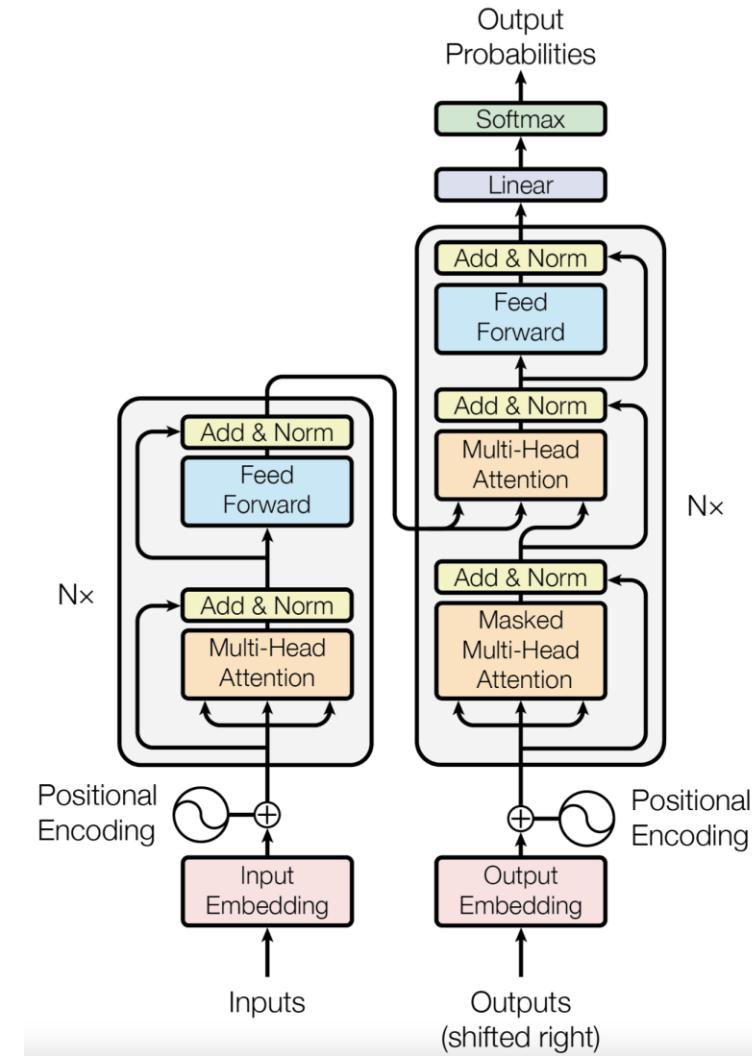
[source](#)

Transformer

attention is all you need: getting rid of RNNs
replaced by multi-headed self-attention (implemented
with matrix multiplications and feed-forward neural
networks)

- allowing for much more parallelization
- allowing for deeper architecture (more parameters)

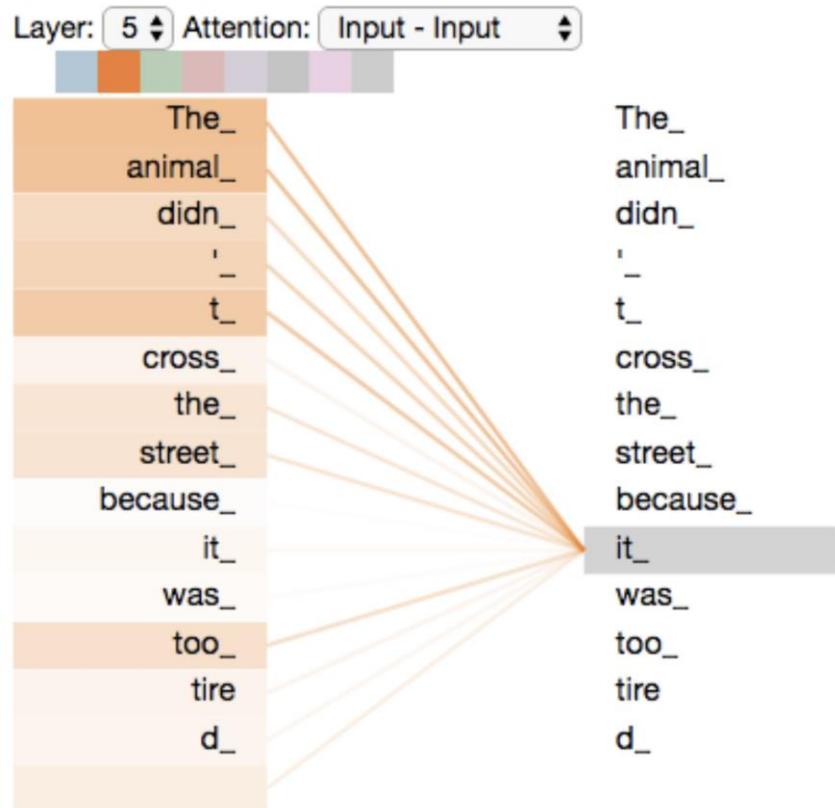
better long-range dependencies thanks to shorter path
lengths in network (less sequential operations)



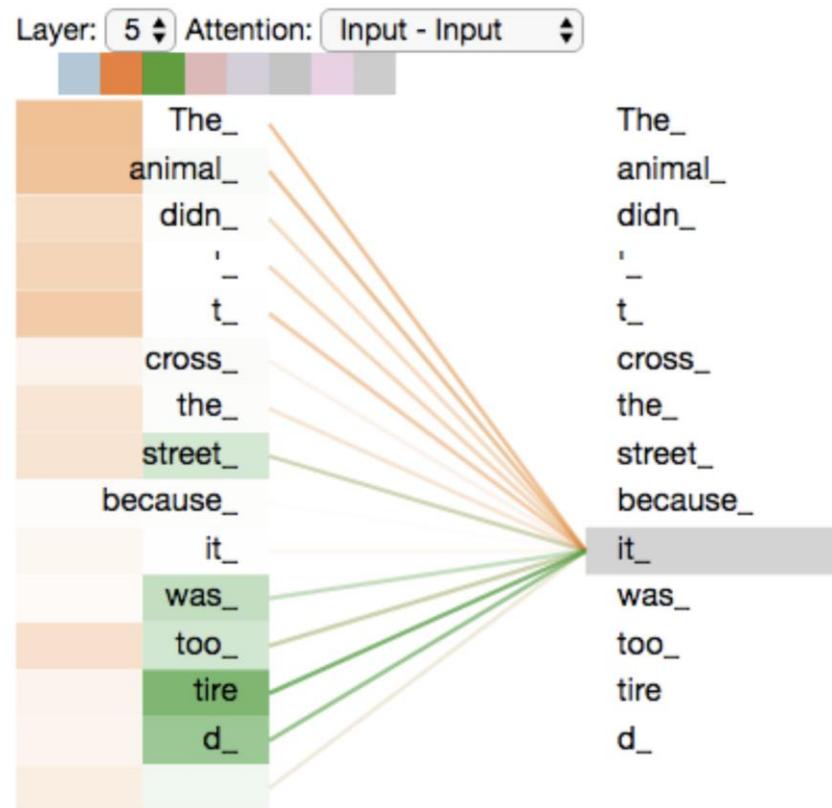
[source](#)

Self-Attention

evaluating other input words in terms of relevance for encoding of given word



multi-head attention: several attention layers running in parallel (considering different aspects of input)



Mechanism: Scaled Dot-Products

3 abstract matrices created from inputs (e.g., word embeddings) by multiplying inputs with 3 different weight matrices

- query Q (embedding dimensions d_k)
- key K (embedding dimensions d_k)
- value V (embedding dimensions d_v)

} context

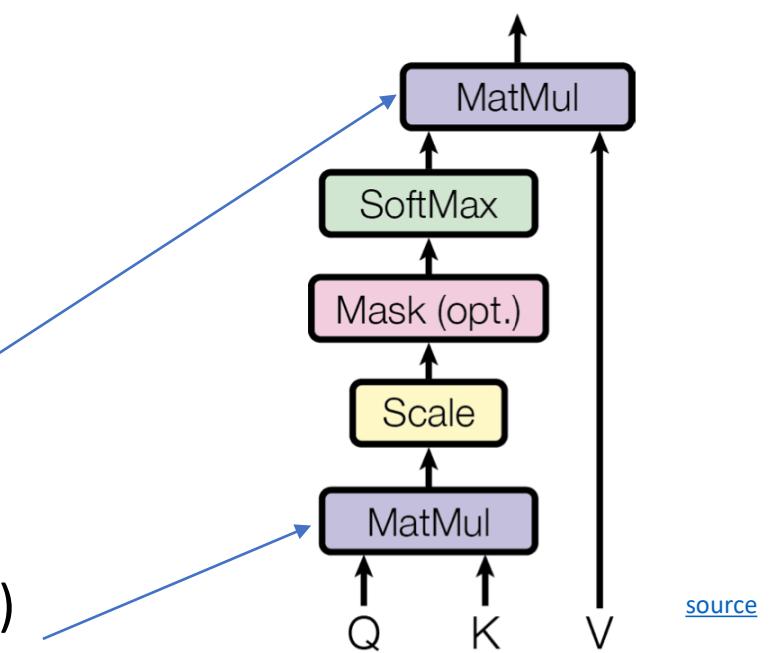
keep values of context words to attend to by multiplication of softmax scores with V

scoring each of the key words (context) with respect to current query word

softmax not scale invariant: largest inputs dominate output for large inputs (more embedding dimensions)

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

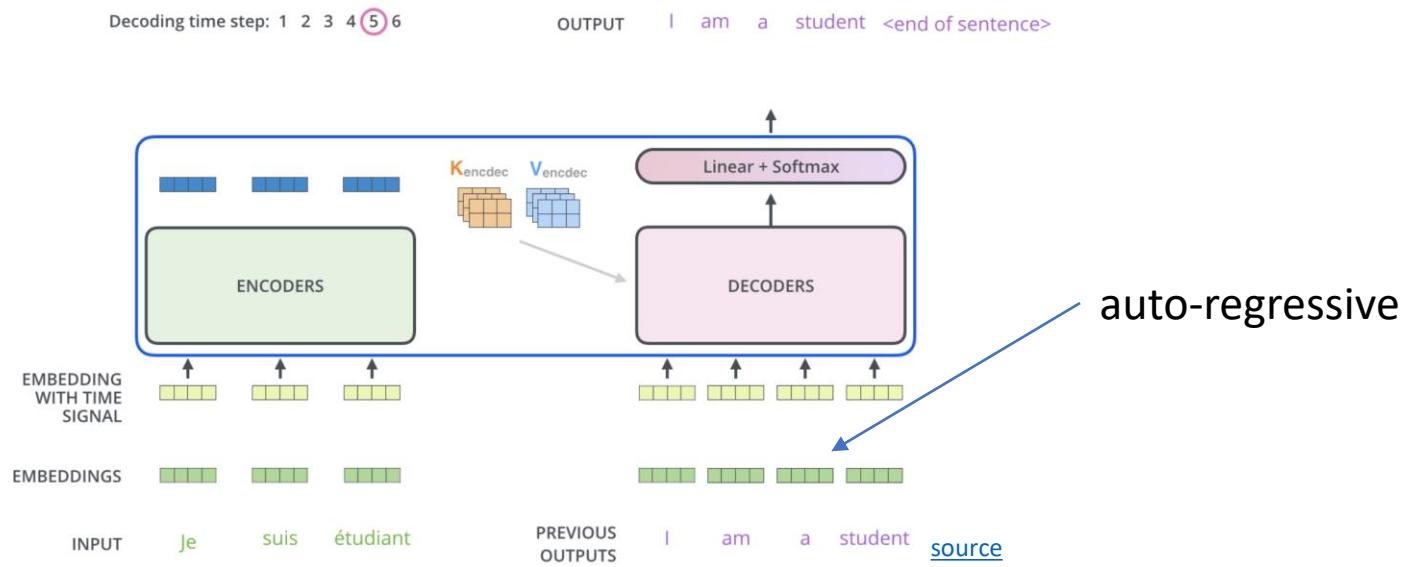
Scaled Dot-Product Attention



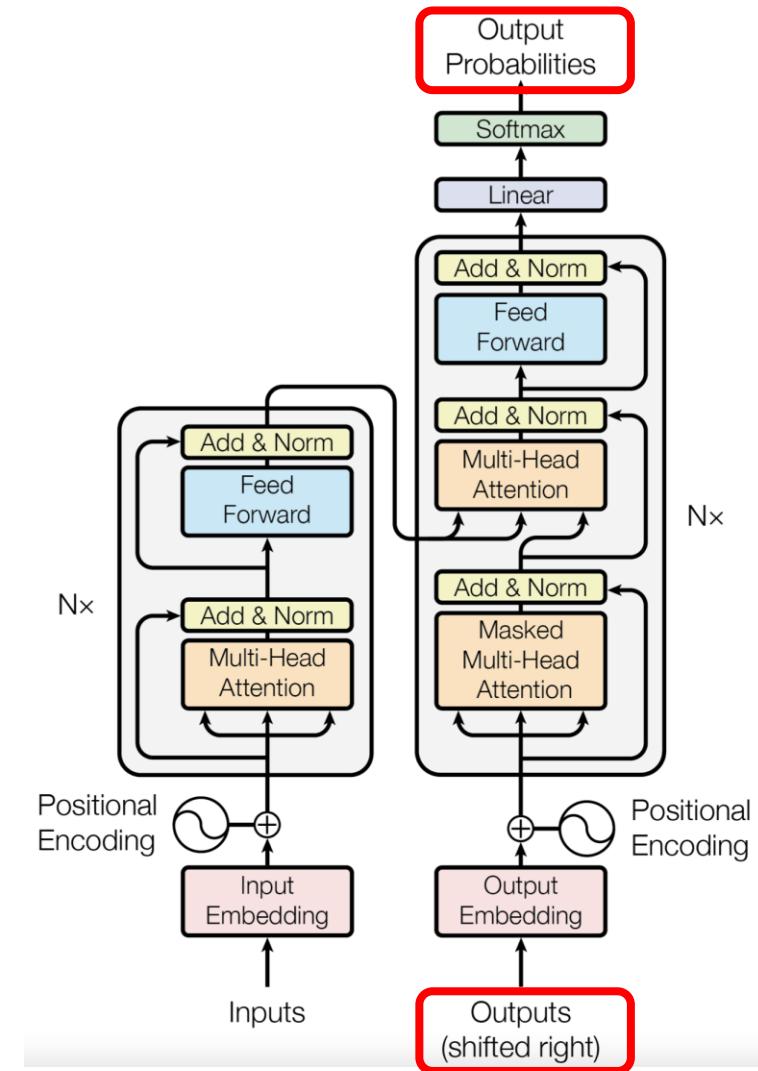
[source](#)

Sequence Completion

for each step/token (iteratively), choose one output token (e.g., greedily picking the one with highest probability or beam search) to add to decoder input sequence → increasing uncertainty



prompt: externally given initial sequence for running start and context on which to build rest of sequence (**prompt engineering**)



[source](#)

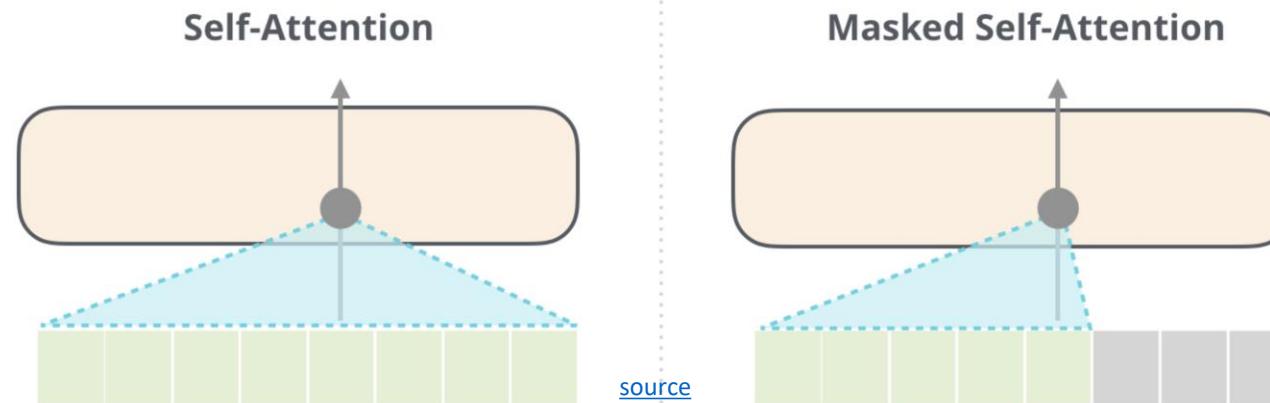
Typical Transformer Architectures for LLMs

encoder-decoder LLMs: sequence-to-sequence, e.g., machine translation

encoder-only LLMs:

- representation learning (and subsequent fine-tuning)
- training: prediction of masked words
- incorporate context of both sides of token

example: Google's [BERT](#)
(Bidirectional Encoder
Representations from
Transformers)



decoder-only LLMs:

- text generation (potentially in-context only), e.g., chat bot
- training: next-word prediction
- output one token at a time (auto-regressive)

example: OpenAI's [GPT](#)
(Generative Pre-trained
Transformer)
[GPT-4 capabilities](#)

Multi-Task Learning of LLMs

compositional nature of deep learning allows transfer learning in a semi-supervised way (also prominent for CNNs in computer vision):

- self-supervised **pre-training** (e.g., next-word prediction) on massive data sets (foundation models like GPT or BERT)
- subsequent supervised **fine-tuning** on specific tasks and (usually much smaller) data sets (by adapting parameters or/and adding layers)

in-context learning as alternative to fine-tuning: only using information fed into LLM via input prompt, no parameter updates (typically decoder-only)

typical prompt: instructions, context (potentially retrieved externally from, e.g., knowledge-base embeddings), query, output indicator

query with (few-shot) or without (zero-shot) providing explicit examples

[prompt engineering](#)

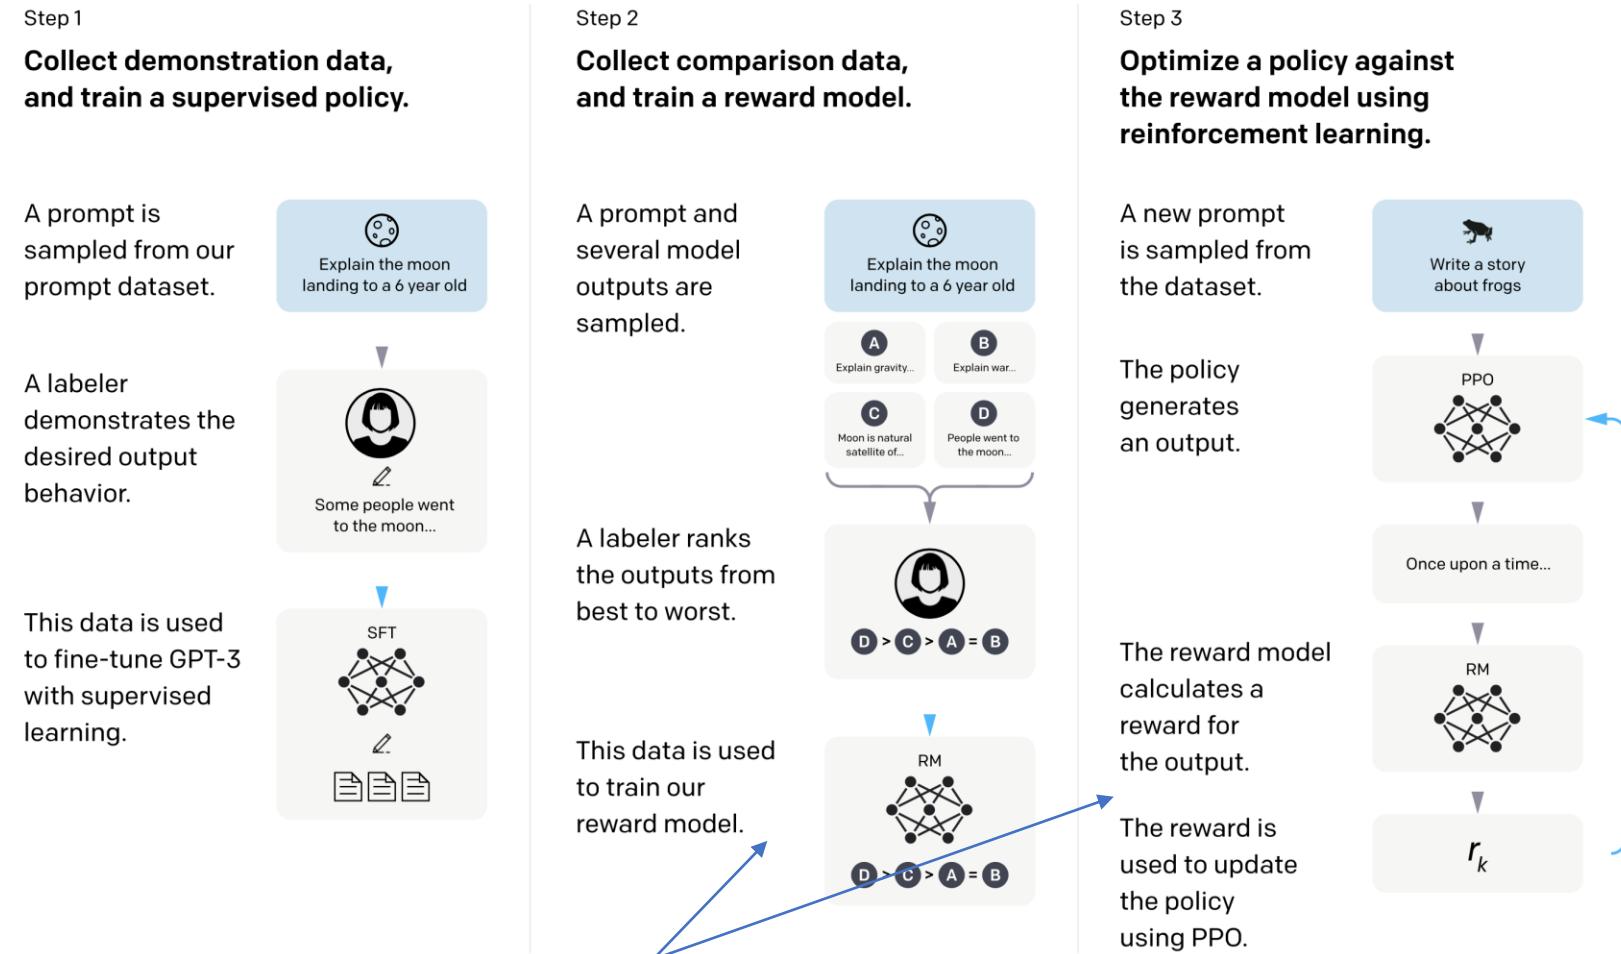
Conversational AI: RL from Human Feedback

example for supporting large language models (transformers) with RL

used in famous ChatGPT

goal: improve alignment with user intentions

→ learn from human preferences



RL looks at reward of text output passages as a whole (rather than token-level loss in supervised learning)

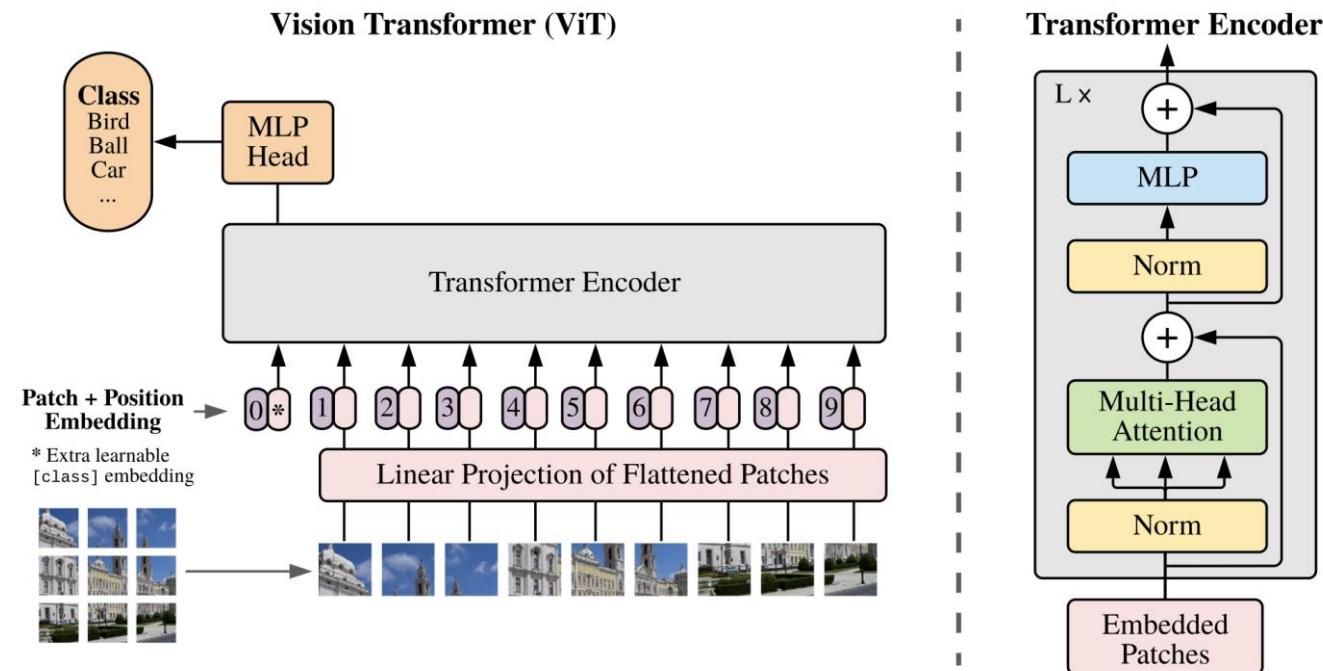
Transformer for Vision

formulation as sequential problem:

- split image into patches and flatten → use as tokens
- produce linear embeddings and add positional embeddings

processing by transformer encoder:

- pre-train with image labels
- fine-tune on specific data set

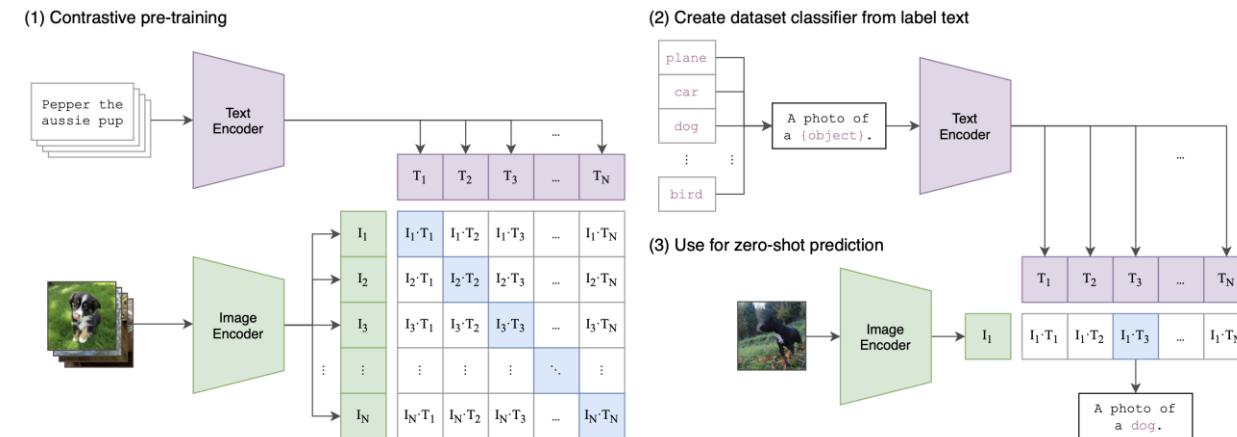


source

Combination of Vision and Text: Multi-Modality

example: [CLIP](#) (Contrastive Language-Image Pre-training)

- learn image representations by predicting which caption goes with which image (pre-training)
- zero-shot transfer (e.g., for object recognition)



multi-modal perception as input for large language models: [KOSMOS-1](#)

multi-purpose (multi-modal and multi-task) models as next generalization step of ML (e.g., Google's [Pathways](#))

transformers good candidate (universal and flexible architecture, little task-specific inductive bias)

Generative Models

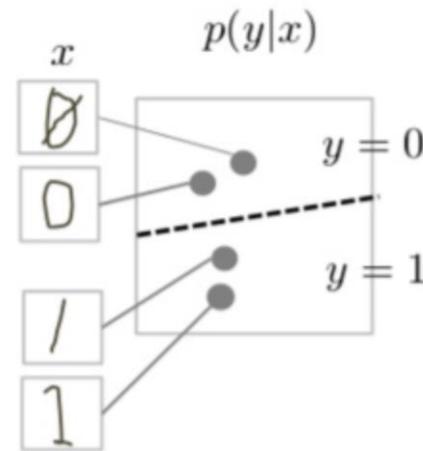
Generative vs Discriminative Models

generative models: predict joint probability $P(Y, X)$ (what allows to create new data samples) or directly generate new data samples

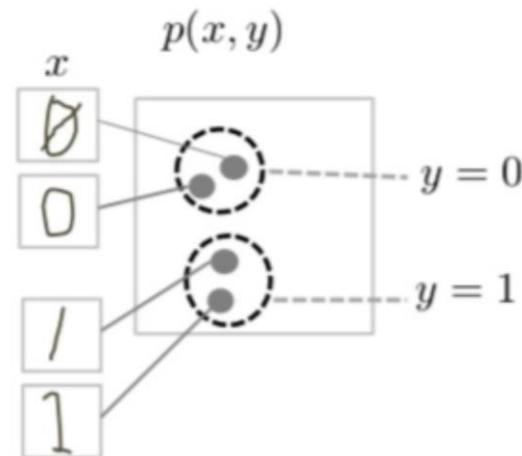
discriminative models: predict conditional probability (or probability distribution for regression) $P(Y|X)$ or directly output (label for classification, real value for regression)

task of generative models more difficult: model full data distribution rather than merely find patterns in inputs to distinguish outputs

discriminative model



generative model



[source](#)

Data Generation

generative models can be used for discriminative tasks (although potentially inferior to direct discriminative methods)

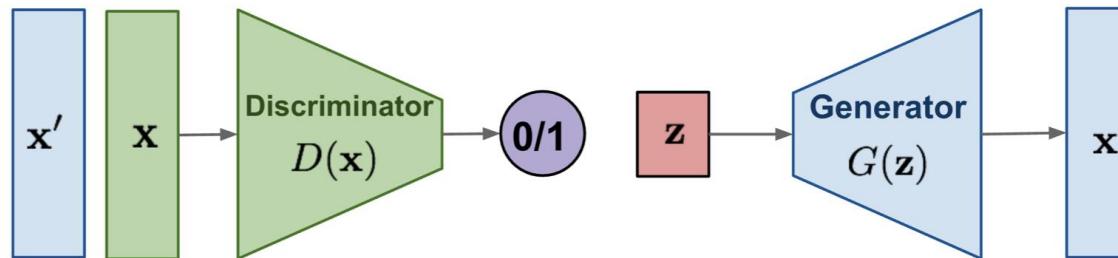
but generative methods do more than discriminative ones: model full data distribution

→ allows generation of new data samples (can be images, text, [video](#), [audio](#), code like SQL or Python, proteins, materials, time series, structured data, ...)

large (auto-regressive) language models examples of generative models

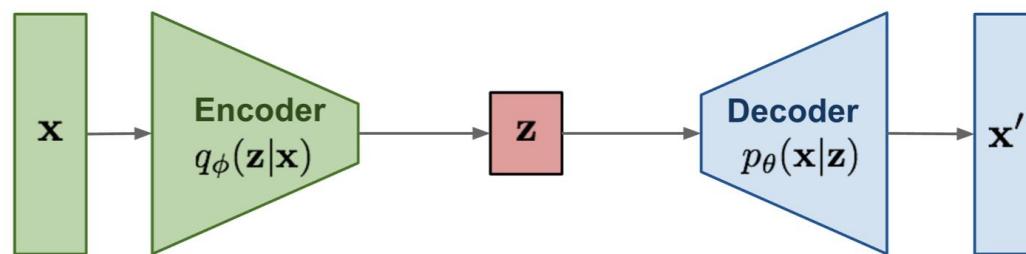
Different Types of Generative Models

GAN: Adversarial training



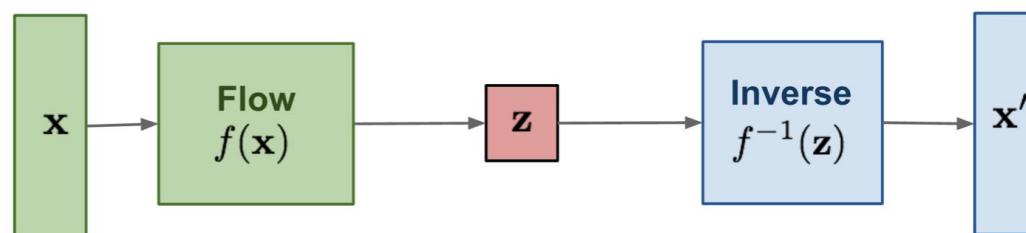
two neural networks playing a zero-sum game

VAE: maximize variational lower bound



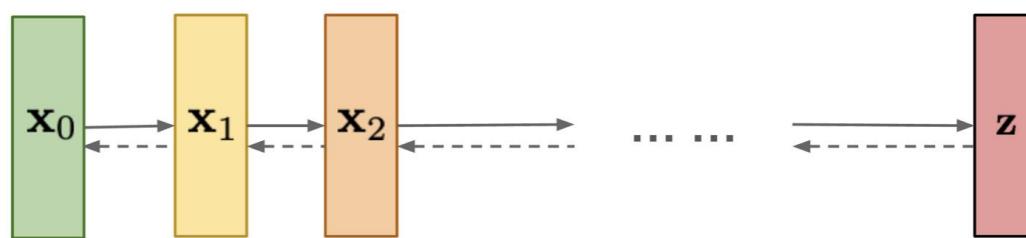
learn variational distribution (not just replicating inputs)

Flow-based models: Invertible transform of distributions



more complex distributions by applying change-of-variable technique (need for specialized architecture)

Diffusion models: Gradually add Gaussian noise and then reverse



chain of denoising autoencoders

Conditioning

tradeoff between diversity (unconditioned) and fidelity (guidance)

as discussed so far, generative methods give no control over what kind of data is generated (limited usability)

→ need for conditional approach (e.g., conditioning on describing text, for example by means of CLIP)

example GANs:

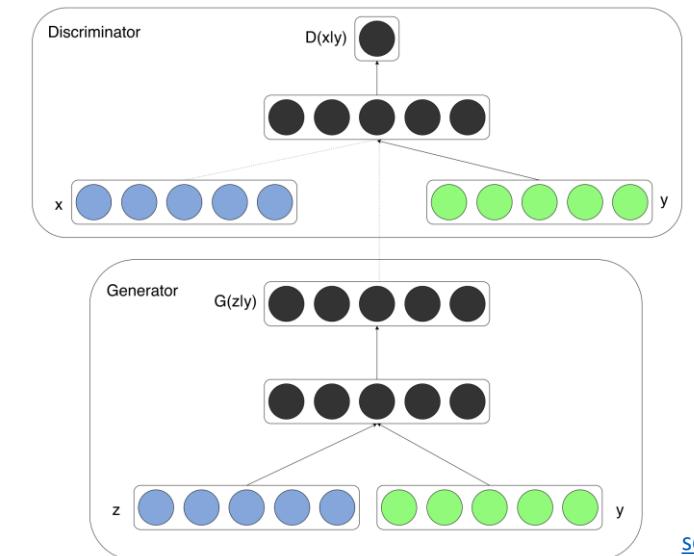
transform usual GAN to conditional model by feeding extra information y (e.g., class labels) as additional input layer into both generator and discriminator

$$L(\mathbf{x}_i) = E_{\mathbf{x} \sim p_r(x)}[\ln D(\mathbf{x}_i | y_i)] + E_{\mathbf{x} \sim p_g(x)}[\ln(1 - D(\mathbf{x}_i | y_i))]$$

G: minimize, D: maximize



guided diffusion: “Pembroke Welsh corgi”



source

Multi-Modal Generative Models

example: generate images from text descriptions

DALL-E (blend of WALL-E and Salvador Dalí): decoder-only transformer auto-regressively modeling text and image tokens as single data stream

TEXT PROMPT

an armchair in the shape of an avocado....

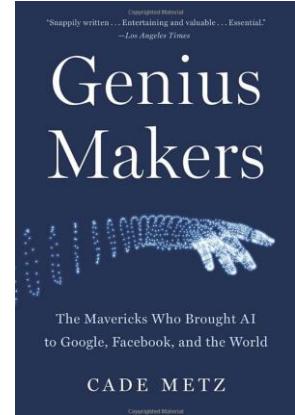
AI-GENERATED IMAGES



[source](#)

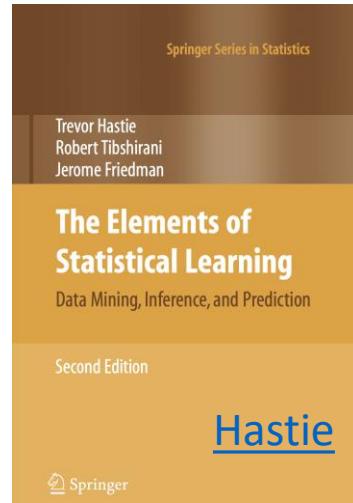
DALL-E 2: image generation conditioned on CLIP image embedding

historical overview
of deep learning:



Literature

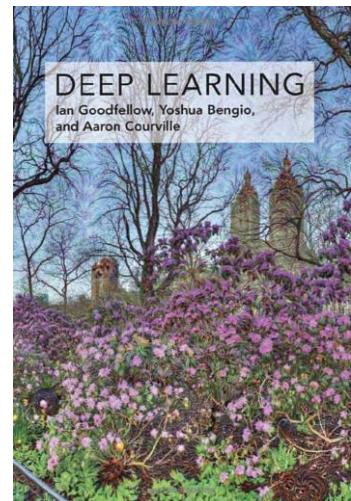
foundations of ML:



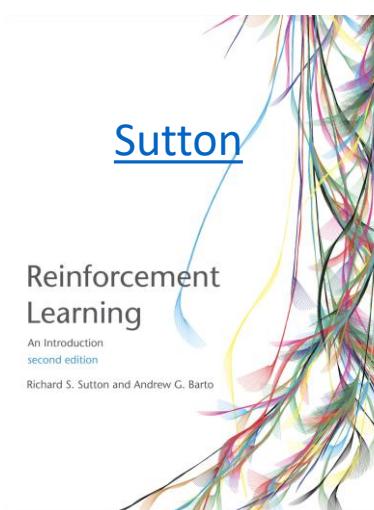
a few seminal papers:

- [back-propagation](#): one of the founding moments of deep learning
- [CNN](#): neural networks work
- [AlexNet](#): deep learning takes over
- [transformer](#): SOTA

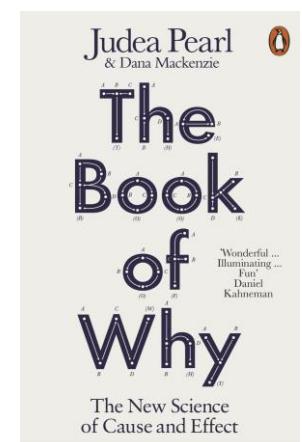
foundations of deep learning:
<https://www.deeplearningbook.org/>



overview of
reinforcement learning:



gentle but genuine
introduction to causality:



Some Philosophical Thoughts

computational theory of mind: mind from matter

just scaling up current methods (e.g., LLMs) enough to achieve general intelligence? or additional methods needed?

agency via goal-based approaches? (is reward enough?)

emergent capabilities of complex systems almost impossible to foresee: emotions or consciousness occur as emergent capabilities?